

Chapter 3: Open Short Path First



CCNP ROUTE: Implementing IP Routing

Cisco | Networking Academy®
Mind Wide Open™



Chapter 3 Objectives

- Basic OSPF Configuration and OSPF Adjacencies
- How OSPF Builds the Routing Table
- Configuration of Summarization and Stub Areas in OSPF
- Configuration of OSPFv3 for IPv6 and IPv4

Establishing OSPF Neighbor Relationships





Establishing OSPF Neighbor Relationships

- Explain why would you choose OSPF over other routing protocols
- Describe basic operation steps with link-state protocols
- Describe area and router types in OSPF
- Explain what the design limitations of OSPF are
- List and describe OSPF message types
- Describe OSPF neighbor relationship over point-to-point link
- Describe OSPF neighbor relationship behavior on MPLS VPN
- Describe OSPF neighbor relationship behavior over L2 MPLS VPN
- List and describe OSPF neighbor states
- List and describe OSPF network types
- Configure passive interfaces



OSPF Features

- OSPF was developed by the Internet Engineering Task Force (IETF) to **overcome the limitations of distance vector routing protocols**.
- One of the main reasons why OSPF is largely deployed in today's enterprise networks is the fact that **it is an open standard**;
- OSPF offers a large level of scalability and **fast convergence**.
- Despite its relatively simple configuration in small and medium-size networks, OSPF implementation and troubleshooting **in large-scale networks can at times be challenging**.



OSPF Features

Independent transport layer protocol

- OSPF works on top of IP and uses protocol number 89.
- It does not rely on the functions of the transport layer protocols TCP or UDP.

Efficient use of updates

- When an OSPF router **first discovers a new neighbor**, it sends a **full update with all known link-state information**.
- **All routers within an OSPF area must have identical and synchronized link-state information in their OSPF link-state databases.**
- When an OSPF network is in a converged state and a new link comes up or a link becomes unavailable, an OSPF router sends **only a partial update to all its neighbors**.
- This update will then be flooded to all OSPF routers within an area.



OSPF Features

Metric

- OSPF uses a metric that is based on the **cumulative costs** of all outgoing interfaces from source to destination. The interface cost is inversely proportional to the interface bandwidth and can be also set up explicitly.

Update destination address

- OSPF uses **multicast** and **unicast**, rather than broadcast, for sending messages.
- The IPv4 multicast addresses used for OSPF are **224.0.0.5** to send information to all OSPF routers and **224.0.0.6** to send information to **DR/BDR** routers.
- The IPv6 multicast addresses are FF02::5 for all OSPFv3 routers and FF02::6 for all DR/BDR routers.
- If the underlying network does not have broadcast capabilities, you must establish OSPF neighbor relationships using a unicast address.
- For IPv6, this address will be a **link-local IPv6 address**.



OSPF Features

VLSM support

- OSPF is a **classless routing protocol**. It supports variable-length subnet masking (VLSM) and discontinuous networks.
- It carries subnet mask information in the routing updates.

Manual route summarization

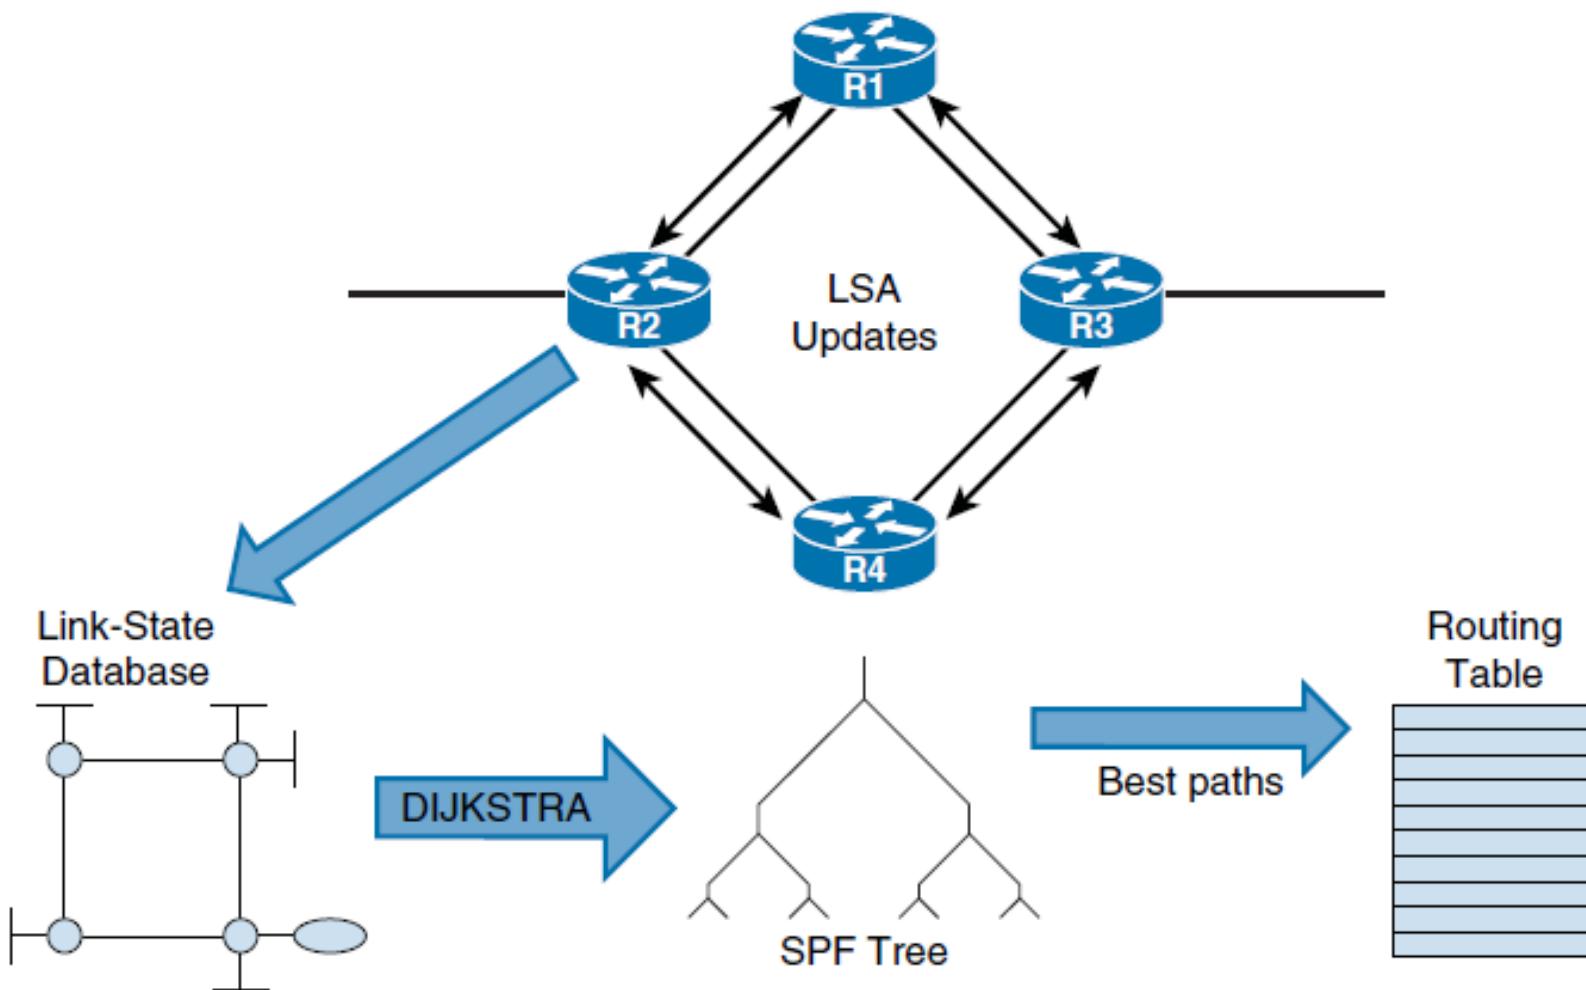
- You can **manually summarize OSPF inter-area routes** at the **Area Border Router (ABR)**, and you have the possibility to summarize OSPF external routes at the **Autonomous System Boundary Router (ASBR)**.
- OSPF does not know the concept of **auto-summarization**.

Authentication

- OSPF supports clear-text, MD5, and SHA authentication.



OSPF Operation Overview





OSPF Operation Overview

1. Establish neighbor adjacencies

- OSPF-enabled routers must form adjacencies with their neighbor before they can share information with that neighbor.
- An OSPF enabled router **sends Hello packets** out all OSPF-enabled interfaces to determine whether neighbors are present on those links.
- If a neighbor is present, the OSPF enabled router attempts to establish a neighbor adjacency with that neighbor.



OSPF Operation Overview

2. Exchange link-state advertisements

- After adjacencies are established, routers then **exchange link-state advertisements (LSAs)**.
- LSAs contain the state and cost of each directly connected link.
- Routers flood their LSAs to adjacent neighbors. Adjacent neighbors receiving the LSA immediately flood the LSA to other directly connected neighbors, **until all routers in the area have all LSAs**.

3. Build the topology table

- After the LSAs are **received**, OSPF-enabled routers **build the topology table** (LSDB) based on the received LSAs.
- This database eventually holds all the information about the topology of the network.
- **It is important that all routers in the area have the same information in their LSDBs.**



OSPF Operation Overview

4. Execute the SPF algorithm

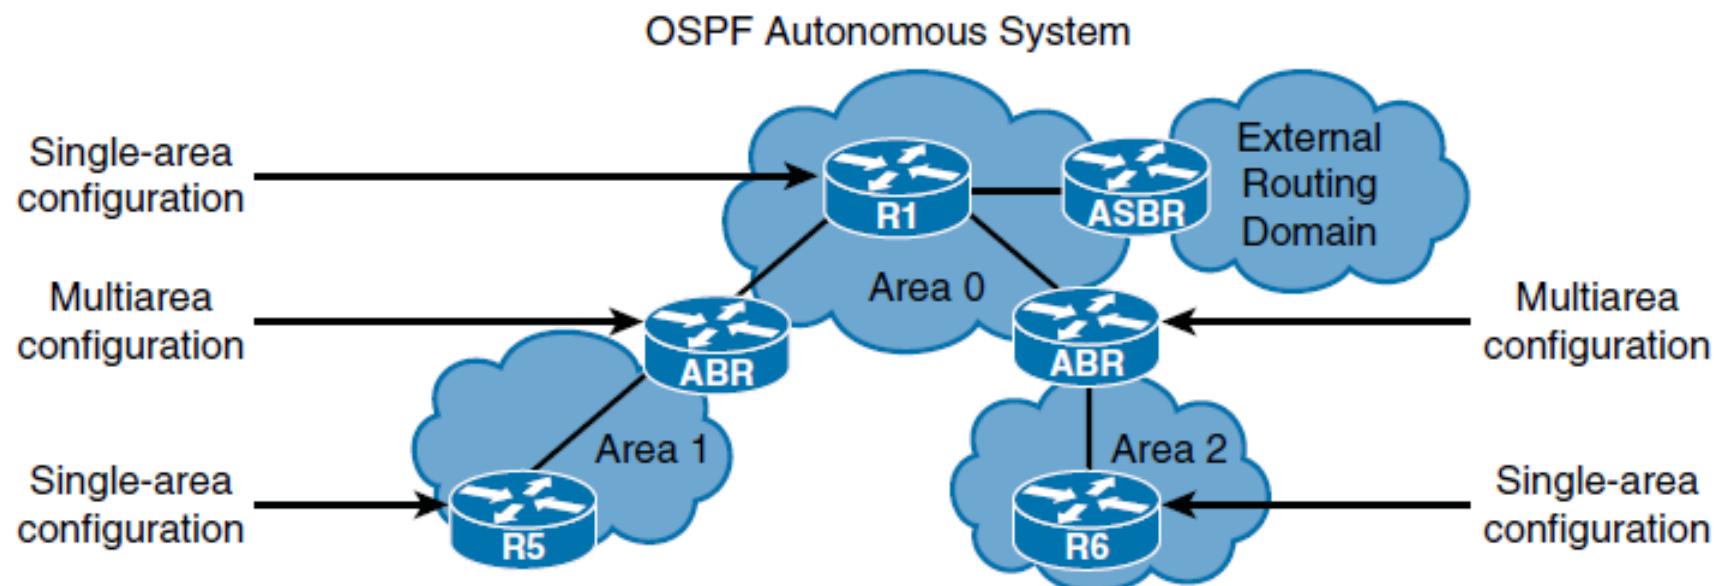
- Routers then execute the SPF algorithm. The SPF algorithm creates the SPF tree.

5. Build the routing table

- From the SPF tree, the best paths are inserted into the routing table. Routing decisions are made based on the entries in the routing table.



Hierarchical Structure of OSPF





Hierarchical Structure of OSPF

OSPF uses a two-layer area hierarchy:

Backbone area, transit area or area 0

- Two principal requirements for the backbone area are that it must connect to all other non-backbone areas and this area must be always contiguous; it is not allowed to have split up the backbone area.
- **Generally, end users are not found within a backbone area.**

Non-backbone area

- The primary function of this area is to connect end users and resources. Non-backbone areas are usually set up according to functional or geographic groupings.
- Traffic between different non-backbone areas must always pass through the backbone area.



Hierarchical Structure of OSPF

- In the multi-area topology there are special commonly OSPF terms:
ABR (Area Border Router)
- A router that has interfaces connected to **at least two different OSPF areas**, including the backbone area. ABRs contain LSDB information for each area, make route calculation for each area and **advertise routing information between areas**.

ASBR (Autonomous System Boundary Router)

- ASBR is a router that has at least one of its interfaces connected to an OSPF area and at least one of its interfaces connected to an external non-OSPF domain.

Internal router

- A router that has all its interfaces connected to only one OSPF area.

Backbone router

- A router that has at least one interface connected to the backbone area.

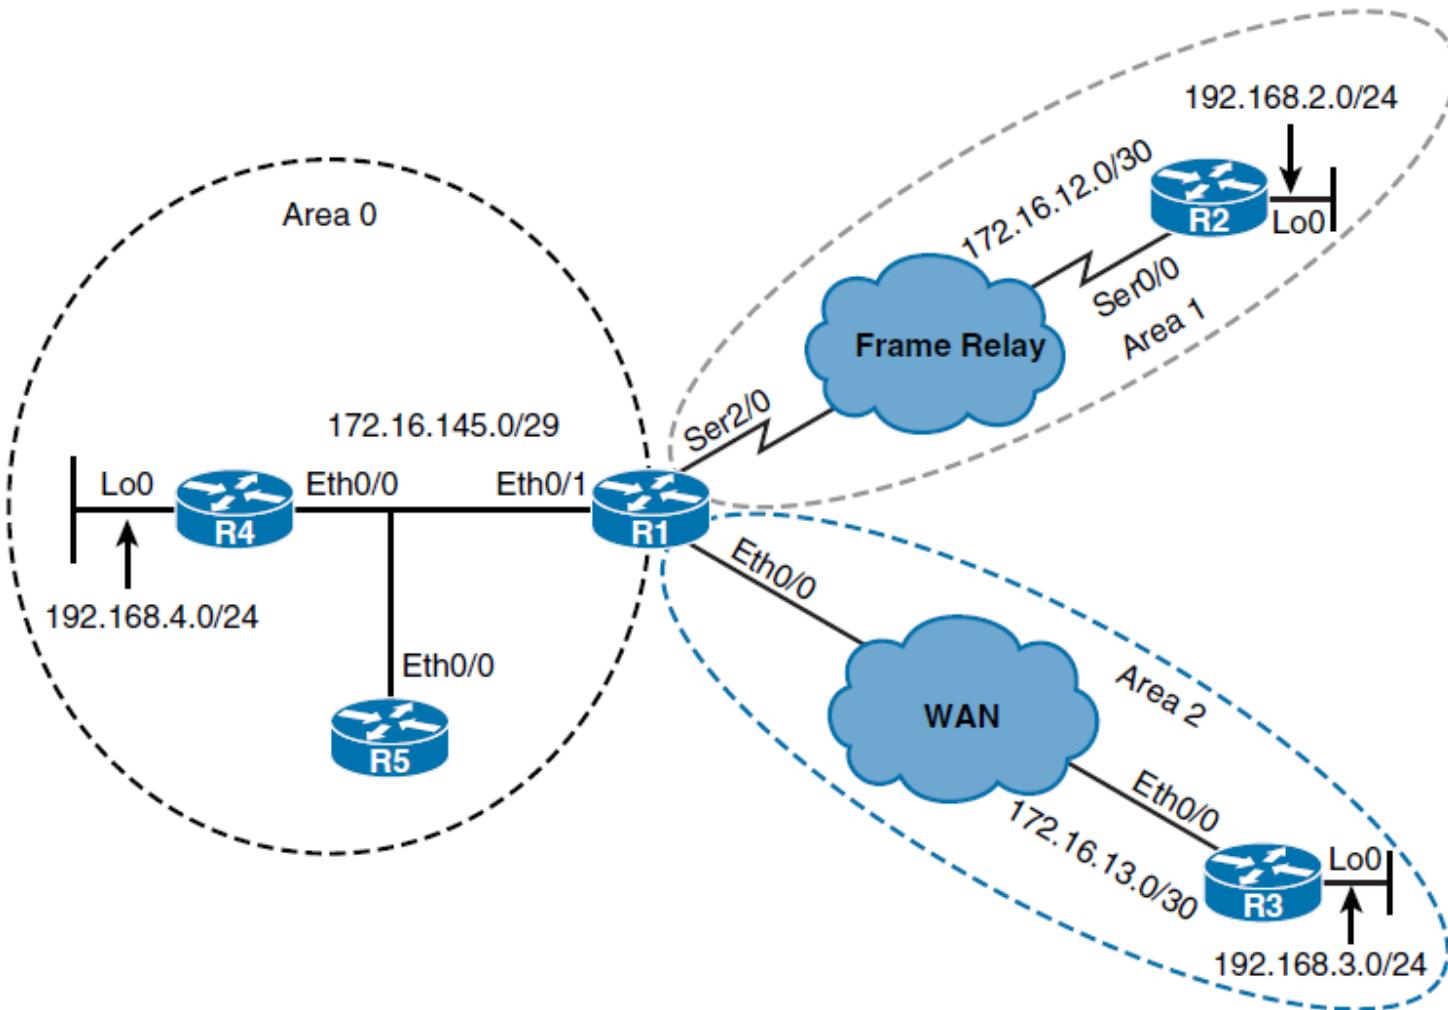


OSPF Packet Types

- **Type 1: Hello packet:** Hello packets are used to discover, build, and maintain OSPF neighbor adjacencies.
- **Type 2: Database Description (DBD) packet:** When the OSPF neighbor adjacency is already established, a DBD packet is used to describe LSDB so that routers can compare whether databases are in sync.
- **Type 3: Link-State Request (LSR) packet :** The router will send an LSR packet to inform OSPF neighbors to send the most recent version of the missing LSAs.
- **Type 4: Link-State Update (LSU) packet :** LSU packets are used for the flooding of LSAs and sending LSA responses to LSR packets.
- **Type 5: Link-State Acknowledgment (LSAck) packet :** LSACKs are used to make flooding of LSAs reliable.



Basic OSPF Configuration





Basic OSPF Configuration

```
R2# configure terminal  
Enter configuration commands, one per line. End with CNTL/Z.  
R2(config)# router ospf 2  
R2(config-router)# network 172.16.12.0 0.0.0.3 area 1  
R2(config-router)# network 192.168.2.0 0.0.0.255 area 1
```

```
R3# configure terminal  
Enter configuration commands, one per line. End with CNTL/Z.  
R3(config)# router ospf 3  
R3(config-router)# network 172.16.13.0 0.0.0.3 area 2  
R3(config-router)# network 192.168.3.0 0.0.0.255 area 2
```

- To enable the OSPF process on the router, use the **router ospf process-id** command.
- Process ID numbers between neighbors do not need to match for the routers to establish an OSPF adjacency.



Basic OSPF Configuration

```
R2# configure terminal  
Enter configuration commands, one per line. End with CNTL/Z.  
R2(config)# router ospf 2  
R2(config-router)# network 172.16.12.0 0.0.0.3 area 1  
R2(config-router)# network 192.168.2.0 0.0.0.255 area 1
```

```
R3# configure terminal  
Enter configuration commands, one per line. End with CNTL/Z.  
R3(config)# router ospf 3  
R3(config-router)# network 172.16.13.0 0.0.0.3 area 2  
R3(config-router)# network 192.168.3.0 0.0.0.255 area 2
```

- To establish OSPF full adjacency, two neighbor routers must be in the same area.
- Any individual interface can only be attached to a single area.



Configuration of OSPF Router IDs

- To choose the OSPF router ID at the time of OSPF process initialization, the router uses the following criteria:
 - Use the router ID specified in the **router-id ip-address** command.
 - Use the highest IPv4 address of all active loopback interfaces on the router.
 - Use the highest IPv4 address among all active nonloopback interfaces.
- At least one primary IPv4 address on an interface in the up/up state must be configured for a router to be able to choose router ID; otherwise, an error message is logged, and the OSPF process does not start.

```
R2(config-router)# router-id 2.2.2.2
% OSPF: Reload or use "clear ip ospf process" command, for this to take effect
```

```
R3(config-router)# router-id 3.3.3.3
% OSPF: Reload or use "clear ip ospf process" command, for this to take effect
```



Clearing the OSPF Processes

- OSPF routing process can be cleared for the manually configured router ID to take effect.

```
R2# clear ip ospf process
Reset ALL OSPF processes? [no]: yes
R2#
*Nov 24 08:37:24.679: %OSPF-5-ADJCHG: Process 2, Nbr 1.1.1.1 on Serial0/0 from
FULL to DOWN, Neighbor Down: Interface down or detached
R2#
*Nov 24 08:39:24.734: %OSPF-5-ADJCHG: Process 2, Nbr 1.1.1.1 on Serial0/0 from
LOADING to FULL, Loading Done
```

```
R3# clear ip ospf 3 process
Reset OSPF process 3? [no]: yes
R3#
*Nov 24 09:06:00.275: %OSPF-5-ADJCHG: Process 3, Nbr 1.1.1.1 on Ethernet0/0 from
FULL to DOWN, Neighbor Down: Interface down or detached
R3#
*Nov 24 09:06:40.284: %OSPF-5-ADJCHG: Process 3, Nbr 1.1.1.1 on Ethernet0/0 from
LOADING to FULL, Loading Done
```



OSPF Router ID Criteria

1. Use the router ID specified in the **router-id ip-address** command.
 - You can configure an arbitrary value in the IPv4 address format, but this value must be unique.
 - If the IPv4 address specified with the **router-id** command overlaps with the router ID of another already-active OSPF process, the **router-id** command fails.
2. Use the highest IPv4 address of all active loopback interfaces on the router.
3. Use the highest IPv4 address among all active nonloopback interfaces.



Verifying the Router IDs

```
R2# show ip protocols
*** IP Routing is NSF aware ***

Routing Protocol is "ospf 2"
  Outgoing update filter list for all interfaces is not set
  Incoming update filter list for all interfaces is not set
  Router ID 2.2.2.2
  Number of areas in this router is 1. 1 normal 0 stub 0 nssa
  Maximum path: 4
  Routing for Networks:
    172.16.12.0 0.0.0.3 area 1
    192.168.2.0 0.0.0.255 area 1
  Routing Information Sources:
    Gateway          Distance      Last Update
    1.1.1.1           110          00:02:55
  Distance: (default is 110)
```

```
R3# show ip protocols | include ID
  Router ID 3.3.3.3
```



Verifying OSPF Neighborships

```
R2# show ip ospf neighbor
```

Neighbor ID	Pri	State	Dead Time	Address	Interface
1.1.1.1	1	FULL/DR	00:01:57	172.16.12.1	Serial0/0

```
R3# show ip ospf neighbor
```

Neighbor ID	Pri	State	Dead Time	Address	Interface
1.1.1.1	1	FULL/DR	00:00:39	172.16.13.1	Ethernet0/0

- **Neighbor ID** : Represents neighbor router ID.
- **Priority** : Priority on the neighbor interface used for the DR/BDR election.
- **State** : OSPF neighbor establishment process
- **Dead Time** : Represents value of the dead timer
- **Address** : Primary IPv4 address of the neighbor router.
- **Interface** : Local interface over which an OSPF neighbor relationship is established.



Verifying the OSPF-Enabled Interfaces

```
R2# show ip ospf interface
Loopback0 is up, line protocol is up
    Internet Address 192.168.2.1/24, Area 1, Attached via Network Statement
        Process ID 2, Router ID 2.2.2.2, Network Type LOOPBACK, Cost: 1
<Output omitted>
Serial0/0 is up, line protocol is up
    Internet Address 172.16.12.2/30, Area 1, Attached via Network Statement
        Process ID 2, Router ID 2.2.2.2, Network Type NON_BROADCAST, Cost: 64
<Output omitted>
```

```
R3# show ip ospf interface
Loopback0 is up, line protocol is up
    Internet Address 192.168.3.1/24, Area 2, Attached via Network Statement
        Process ID 3, Router ID 3.3.3.3, Network Type LOOPBACK, Cost: 1
<Output omitted>
Ethernet0/0 is up, line protocol is up
    Internet Address 172.16.13.2/30, Area 2, Attached via Network Statement
        Process ID 3, Router ID 3.3.3.3, Network Type BROADCAST, Cost: 10
<Output omitted>
```



Verifying the OSPF Routes

```
R5# show ip route ospf
```

Gateway of last resort is not set

172.16.0.0/16 is variably subnetted, 4 subnets, 3 masks

```
O IA    172.16.12.0/30 [110/74] via 172.16.145.1, 00:39:00, Ethernet0/0
O IA    172.16.13.0/30 [110/20] via 172.16.145.1, 00:19:29, Ethernet0/0
      192.168.2.0/32 is subnetted, 1 subnets
O IA    192.168.2.1 [110/75] via 172.16.145.1, 00:07:27, Ethernet0/0
      192.168.3.0/32 is subnetted, 1 subnets
O IA    192.168.3.1 [110/21] via 172.16.145.1, 00:08:30, Ethernet0/0
O        192.168.4.0/24 [110/11] via 172.16.145.4, 00:39:10, Ethernet0/0
```



OSPF Routes

```
R5# show ip ospf route

        OSPF Router with ID (5.5.5.5) (Process ID 1)

        Base Topology (MTID 0)

        Area BACKBONE(0)

        Intra-area Route List
*   172.16.145.0/29, Intra, cost 10, area 0, Connected
    via 172.16.145.5, Ethernet0/0
*-> 192.168.4.0/24, Intra, cost 11, area 0
    via 172.16.145.4, Ethernet0/0

        Intra-area Router Path List
i 1.1.1.1 [10] via 172.16.145.1, Ethernet0/0, ABR, Area 0, SPF 2

        Inter-area Route List
*> 192.168.2.1/32, Inter, cost 75, area 0
    via 172.16.145.1, Ethernet0/0
*> 192.168.3.1/32, Inter, cost 21, area 0
    via 172.16.145.1, Ethernet0/0
*> 172.16.12.0/30, Inter, cost 74, area 0
    via 172.16.145.1, Ethernet0/0
*> 172.16.13.0/30, Inter, cost 20, area 0
    via 172.16.145.1, Ethernet0/0
```

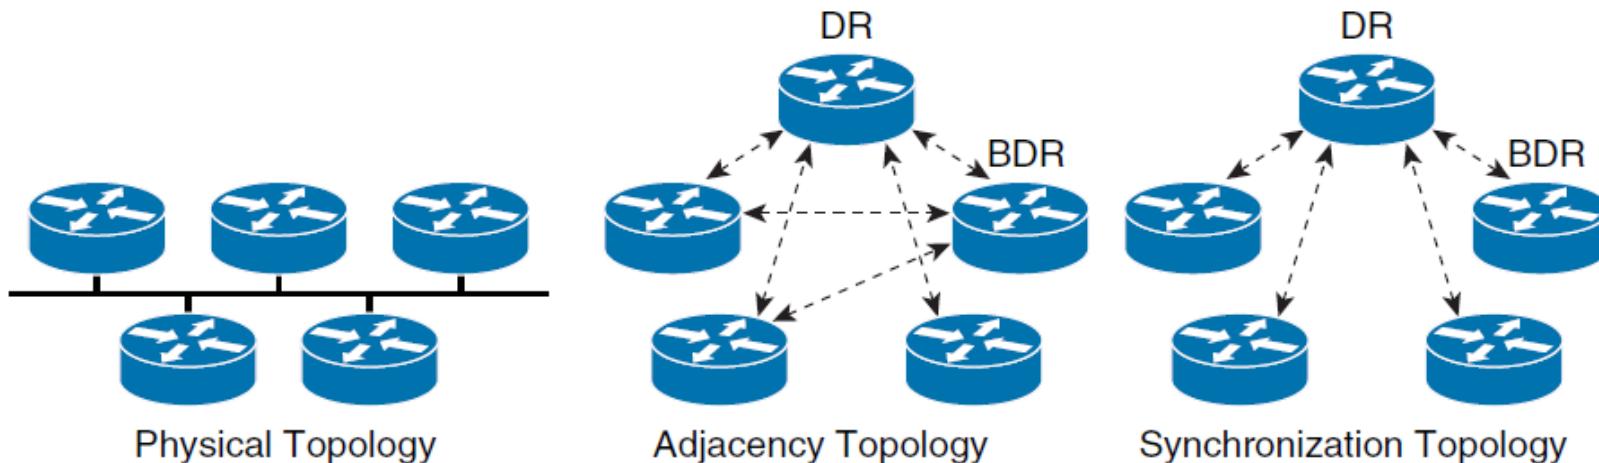


OSPF Adjacency Establishment Steps

1. **Down state** - the initial state of a neighbor conversation that indicates that **no Hello's have been heard from the neighbor**.
2. **Init state** - a router **receives a Hello** from the neighbor but has **not yet seen its own router ID** in the neighbor Hello packet
3. **2-Way state** - When the router sees its own router ID in the Hello packet received from the neighbor
4. **ExStart state** - In the DBD exchange process, the router with the **higher router ID will become master**, and it will be the only router that can increment sequence numbers. With master/slave selection complete, database exchange can start
5. **Exchange state** - To describe the content of the database, one or multiple DBD packets may be exchanged. A router compares the content of its own Database Summary list with the list received from the neighbor, and if there are differences, it adds missing LSAs to the Link State Request list.
6. **Loading state** - LSR packet is sent to the neighbor requesting full content of the missing LSAs from the LS Request list
7. **Full state** - Finally, when neighbors have a complete version of the LSDB, which means that databases on the routers are synchronized and that neighbors are fully adjacent.



Optimizing OSPF Adjacency Behavior



- As the number of routers on the segment grows, the number of OSPF adjacencies increases exponentially. Every router must synchronize its OSPF database with every other router, and in the case of a **large number of routers**, this **leads to inefficiency**.



OSPF Adjacencies on Multiaccess Networks

- The routers on the multiaccess segment elect a designated router (**DR**) and backup designated router (**BDR**), which centralizes communications for all routers connected to the segment.
- The DR and BDR improve network functioning in the following ways:
 - **Reducing routing update traffic:** The DR and BDR act as a central point of contact for link-state information exchange on a multiaccess network; therefore, each router must establish a full adjacency with the DR and the BDR only.
 - **Managing link-state synchronization:** The DR and BDR ensure that the other routers on the network have the same link-state information about the common segment.



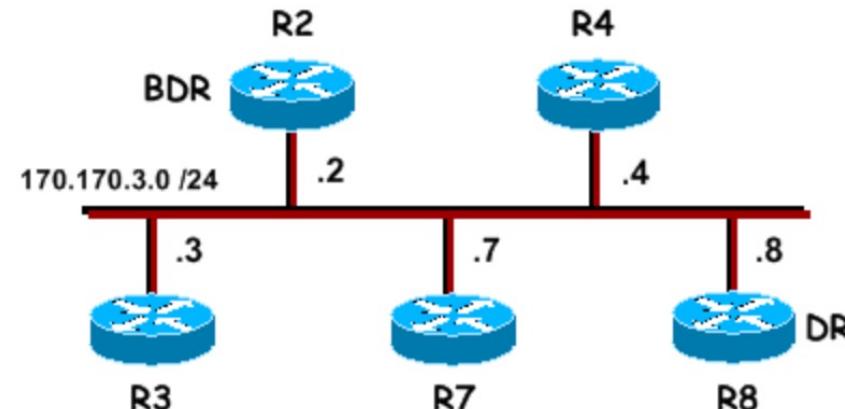
OSPF Adjacencies on Multiaccess Networks

- Only LSAs are sent to the DR/BDR. **The normal routing of packets on the segment will go to the best next-hop router.**
- When the DR is operating, the BDR does not perform any DR functions. Instead, the **BDR receives all the information, but the DR performs the LSA forwarding and LSDB synchronization tasks.**
- The BDR performs the DR tasks **only if the DR fails.**
- When the DR fails, the **BDR automatically becomes the new DR**, and a new BDR election occurs.



Neighbor Status

normal behavior for OSPF



R7# show ip ospf neighbor

Neighbor ID	Pri	State	Dead Time	Address	Interface
170.170.3.4	1	2WAY/DROTHER	00:00:34	170.170.3.4	Ethernet0
170.170.3.3	1	2WAY/DROTHER	00:00:34	170.170.3.3	Ethernet0
170.170.3.8	1	FULL/DR	00:00:32	170.170.3.8	Ethernet0
170.170.3.2	1	FULL/BDR	00:00:39	170.170.3.2	Ethernet0

R8# show ip ospf neighbor

Neighbor ID	Pri	State	Dead Time	Address	Interface
170.170.3.4	1	FULL/DROTHER	00:00:37	170.170.3.4	Ethernet0
170.170.3.3	1	FULL/DROTHER	00:00:37	170.170.3.3	Ethernet0
170.170.3.7	1	FULL/DROTHER	00:00:38	170.170.3.7	Ethernet0
170.170.3.2	1	FULL/BDR	00:00:32	170.170.3.2	Ethernet0



Neighbor Status

```
R1# show ip ospf neighbor
```

Neighbor ID	Pri	State	Dead Time	Address	Interface
4.4.4.4	1	FULL/BDR	00:00:37	172.16.145.4	Ethernet0/1
5.5.5.5	1	FULL/DR	00:00:39	172.16.145.5	Ethernet0/1
2.2.2.2	1	FULL/DR	00:01:53	172.16.12.2	Serial2/0
3.3.3.3	1	FULL/DR	00:00:35	172.16.13.2	Ethernet0/0

```
R4# show ip ospf neighbor
```

Neighbor ID	Pri	State	Dead Time	Address	Interface
1.1.1.1	1	FULL/DROTHER	00:00:39	172.16.145.1	Ethernet0/0
5.5.5.5	1	FULL/DR	00:00:39	172.16.145.5	Ethernet0/0

```
R5# show ip ospf neighbor
```

Neighbor ID	Pri	State	Dead Time	Address	Interface
1.1.1.1	1	FULL/DROTHER	00:00:39	172.16.145.1	Ethernet0/0
4.4.4.4	1	FULL/BDR	00:00:35	172.16.145.4	Ethernet0/0



DR/BDR Election Process

- The DR/BDR election process takes place on **broadcast and NBMA networks**.
- Routers send OSPF Hello packets to discover which OSPF neighbors are active on the common Ethernet segment.
- Once OSPF neighbors are in the 2-Way state, the DR/BDR election process begins.
- The OSPF Hello packet contains three specific fields used for the DR/BDR election:
 - Designated Router, Backup Designated Router, and Router Priority.
- From all routers listed in this fields, the **router with the highest priority becomes the DR**, and the one with the next highest priority becomes the BDR.
- If the priority values are equal, the **router with the highest OSPF router ID becomes the DR**, and the one with the next highest OSPF router ID becomes the BDR.



DR/BDR Election Process

- On the multiaccess broadcast networks, routers use multicast destination **IPv4 address 224.0.0.6** to communicate with the DR, and the DR uses multicast destination **IPv4 address 224.0.0.5** to communicate with all other non-DR routers
- On NBMA networks, the DR and adjacent routers communicate using **unicast addresses**.
- The DR/BDR election process not only occurs when the network first becomes active but also when the DR becomes unavailable.
- In this case, the BDR will immediately become the DR, and the election of the new BDR starts.
- On the multi-access segment, it is normal behavior that the router in **DROTHER** status is **FULLy** adjacent with DR/BDR and in **2-WAY** state with **all other DROTHER routers** present on the segment.



Using OSPF Priority in the DR/BDR Election

Every broadcast and NBMA OSPF-enabled interface is assigned a priority value between 0 and 255. By default, the OSPF interface priority is 1 and can be manually changed by using the **ip ospf priority** interface command.

- The router with the **highest priority** value is elected as the **DR**.
- The router with the **second-highest priority** value is the **BDR**.
- In case of a tie where two routers have the same priority value, router ID is used as the tiebreaker. The router with the highest router ID becomes the DR. The router with the second-highest router ID becomes the BDR.
- A router with a **priority** that is set to **0 cannot become the DR or BDR**. A router that is not the DR or BDR is called a DROTHER.



Configuring the OSPF Priority on an Interface

```
R1(config)# interface ethernet 0/1
R1(config-if)# ip ospf priority 100
```

```
R4# clear ip ospf process
Reset ALL OSPF processes? [no] : yes
*Dec 10 13:08:48.610: %OSPF-5-ADJCHG: Process 1, Nbr 1.1.1.1 on Ethernet0/0 from
FULL to DOWN, Neighbor Down: Interface down or detached
*Dec 10 13:08:48.610: %OSPF-5-ADJCHG: Process 1, Nbr 5.5.5.5 on Ethernet0/0 from
FULL to DOWN, Neighbor Down: Interface down or detached
*Dec 10 13:09:01.294: %OSPF-5-ADJCHG: Process 1, Nbr 1.1.1.1 on Ethernet0/0 from
LOADING to FULL, Loading Done
*Dec 10 13:09:04.159: %OSPF-5-ADJCHG: Process 1, Nbr 5.5.5.5 on Ethernet0/0 from
LOADING to FULL, Loading Done
```



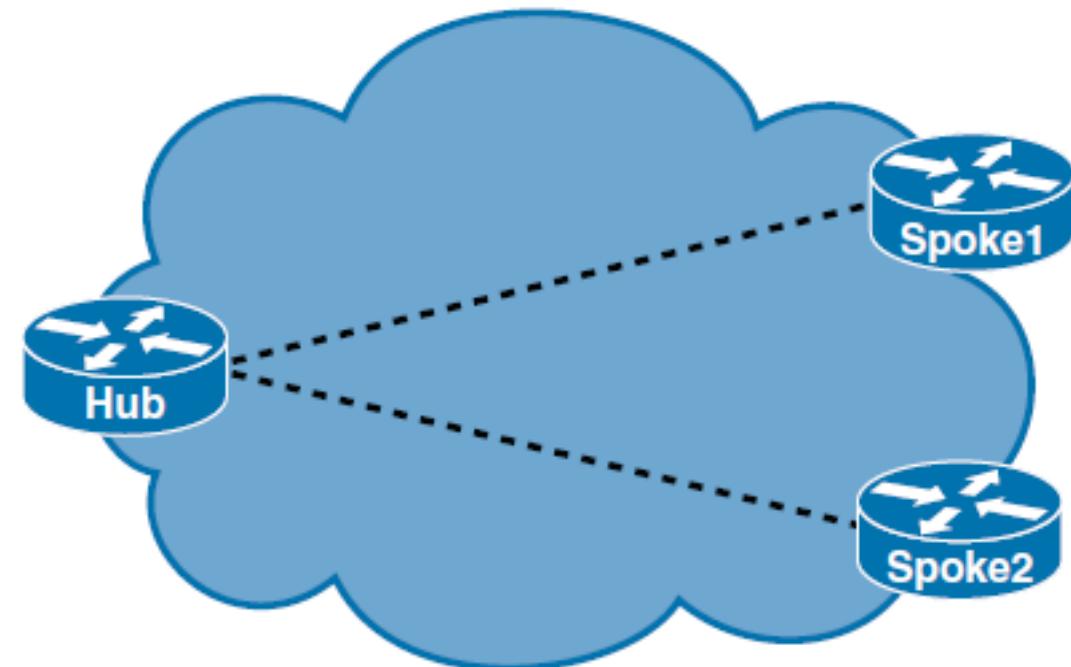
Verifying OSPF Neighbor Status

```
R1# show ip ospf interface ethernet 0/1
Ethernet0/1 is up, line protocol is up
  Internet Address 172.16.145.1/29, Area 0, Attached via Network Statement
  Process ID 1, Router ID 1.1.1.1, Network Type BROADCAST, Cost: 10
  Topology-MTID      Cost      Disabled      Shutdown      Topology Name
    0                  10        no            no            Base
  Transmit Delay is 1 sec, State DR, Priority 100
  Designated Router (ID) 1.1.1.1, Interface address 172.16.145.1
  Backup Designated router (ID) 5.5.5.5, Interface address 172.16.145.5
  Timer intervals configured, Hello 10, Dead 40, Wait 40, Retransmit 5
    oob-resync timeout 40
  Hello due in 00:00:06
  Supports Link-local Signaling (LLS)
  Cisco NSF helper support enabled
  IETF NSF helper support enabled
  Index 1/3, flood queue length 0
  Next 0x0(0)/0x0(0)
  Last flood scan length is 1, maximum is 5
  Last flood scan time is 0 msec, maximum is 1 msec
  Neighbor Count is 2, Adjacent neighbor count is 2
    Adjacent with neighbor 4.4.4.4
    Adjacent with neighbor 5.5.5.5 (Backup Designated Router)
  Suppress hello for 0 neighbor(s)
```



OSPF Behavior in NBMA Hub-and-Spoke Topology

- OSPF treats NBMA environments like any other broadcast media environment, such as Ethernet; however, NBMA clouds are usually built as hub-and-spoke topologies using private virtual circuits (PVCs) or switched virtual circuits (SVCs)





OSPF Behavior in NBMA Hub-and-Spoke Topology

- In these cases, the physical topology does not provide multiaccess capability, on which OSPF relies. In a hub-and-spoke NBMA environment, you will **need to have the hub router acting as the DR and spoke routers acting as the DROTHER routers**. On the spoke router interfaces, you want to configure an OSPF priority value of 0 so that the spoke routers never participate in the DR election.

```
R4 (config)# interface ethernet 0/0
R4 (config-if)# ip ospf priority 0
```

```
R5 (config)# interface ethernet 0/0
R5 (config-if)# ip ospf priority 0
```



The Importance of MTU

- If neighbors have a mismatched IPv4 MTU configured, they will not be able to form full OSPF adjacency.
- Mismatched neighbors will stay in ExStart state

```
R3(config)# interface ethernet 0/0
R3(config-if)# ip mtu 1400
```

- By default, the IPv6 MTU must also match between OSPFv3 neighbors. However, you can override this by using the **ospfv3 mtu-ignore** interface command.



Manipulating OSPF Timers

- Similar to EIGRP, OSPF uses two timers to check neighbor reachability: the hello and dead intervals.
- OSPF requires that both **hello and dead** timers **be identical for all routers on the segment** to become OSPF neighbors.
- The default value of the OSPF hello timer on multi-access broadcast and point-to-point links is **10 seconds**, and is 30 seconds on all other network types, including NBMA.
- When you configure the hello interval, the default value of the dead interval is automatically adjusted to **4 times the hello interval**.
- For broadcast and point-to-point links, it is 40 seconds, and for all other OSPF network types, it is 120 seconds.



Examining the Hello/Dead Timers

```
R1# show ip ospf interface ethernet 0/1
Ethernet0/1 is up, line protocol is up
  Internet Address 172.16.145.1/29, Area 0, Attached via Network Statement
  Process ID 1, Router ID 1.1.1.1, Network Type BROADCAST, Cost: 10
    Topology-MTID      Cost      Disabled      Shutdown      Topology Name
      0                10        no            no            Base
    Transmit Delay is 1 sec, State DROTHER, Priority 1
    Designated Router (ID) 5.5.5.5, Interface address 172.16.145.5
    Backup Designated router (ID) 4.4.4.4, Interface address 172.16.145.4
    Timer intervals configured, Hello 10, Dead 40, Wait 40, Retransmit 5
<Output omitted>
```

```
R1# show ip ospf interface serial 2/0
Serial2/0 is up, line protocol is up
  Internet Address 172.16.12.1/30, Area 1, Attached via Network Statement
  Process ID 1, Router ID 1.1.1.1, Network Type NON_BROADCAST, Cost: 64
    Topology-MTID      Cost      Disabled      Shutdown      Topology Name
      0                64        no            no            Base
    Transmit Delay is 1 sec, State BDR, Priority 1
    Designated Router (ID) 2.2.2.2, Interface address 172.16.12.2
    Backup Designated router (ID) 1.1.1.1, Interface address 172.16.12.1
    Timer intervals configured, Hello 30, Dead 120, Wait 120, Retransmit 5
<Output omitted>
```



Modifying the Hello and Dead Intervals

```
R1(config)# interface serial 2/0
R1(config-if)# ip ospf hello-interval 8
R1(config-if)# ip ospf dead-interval 30
*Jan 20 13:17:34.441: %OSPF-5-ADJCHG: Process 1, Nbr 2.2.2.2 on Serial2/0 from
FULL to DOWN, Neighbor Down: Dead timer expired
```



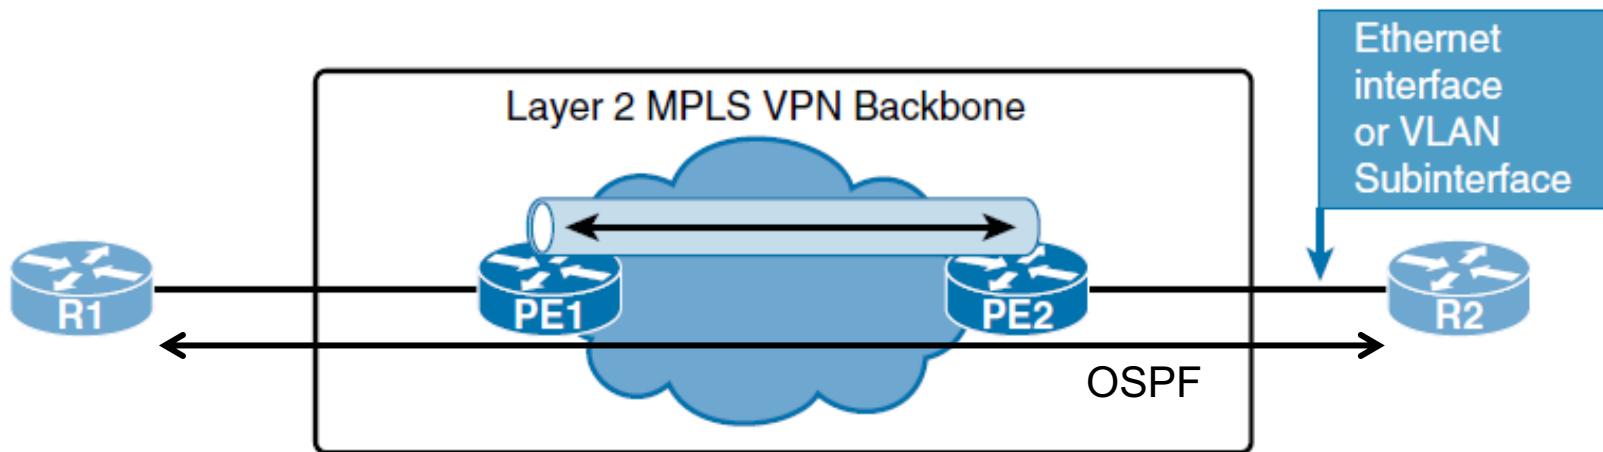
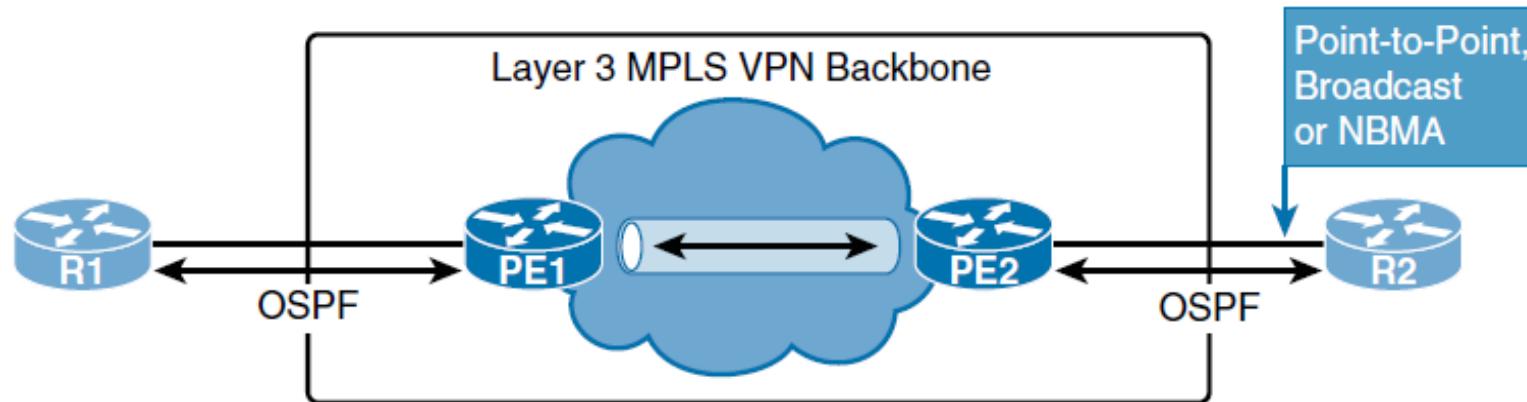
OSPF Relationship over Point-to-Point Links

- The routers dynamically detects its neighboring routers by multicasting its Hello packets to all OSPF routers, using the 224.0.0.5 address.
- **No DR or BDR election is performed**
- The default OSPF hello and dead timers on point-to-point links are 10 seconds and 40 seconds, respectively.



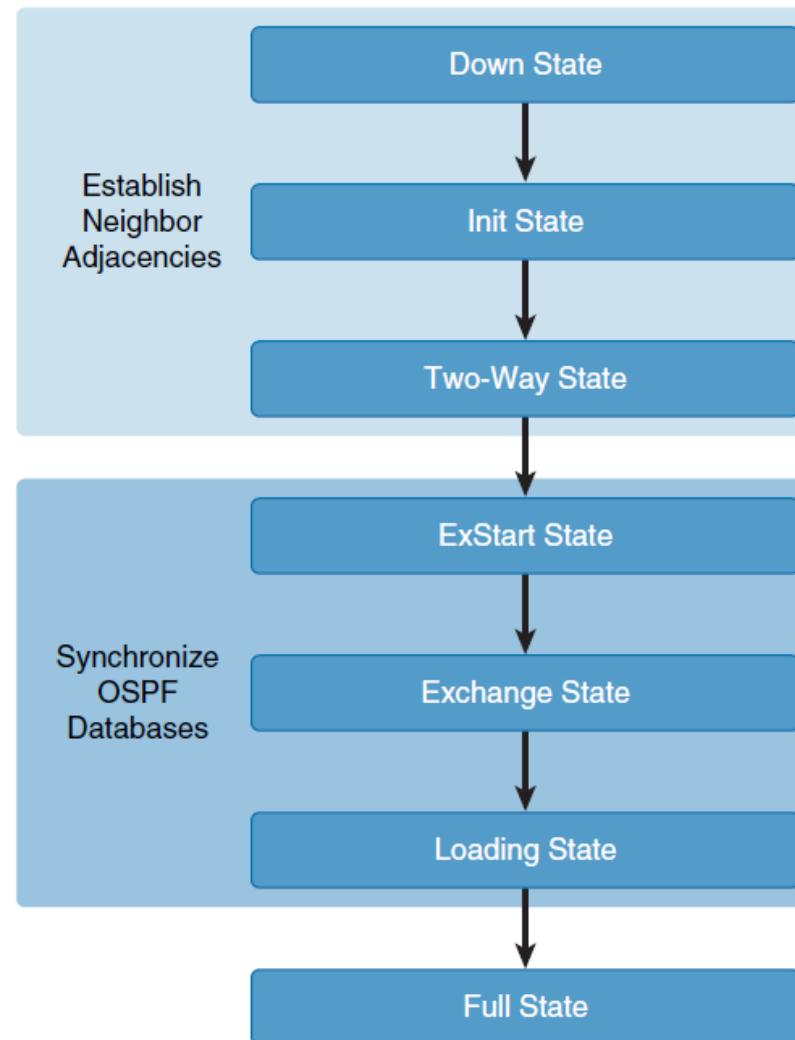


OSPF Neighbor Relationship over MPLS





OSPF Neighbor States





- **Down:** No information has been received on the segment.
- **Init:** The interface has detected a Hello packet coming from a neighbor, but bidirectional
- **2-Way:** There is bidirectional communication with a neighbor. The router **has seen itself in the Hello packets** coming from a neighbor
- **ExStart:** Routers are trying to establish the **initial sequence number** that is going to be used in the information exchange packets.
- **Exchange:** Routers will describe their entire LSDB by sending database description (DBD) packets.
- **Loading:** In this state, routers are finalizing the information exchange.
- **Full:** In this state, adjacency is complete.



OSPF Network Types

OSPF Network Type	Uses DR/BDR	Default Hello Interval (sec)	Dynamic Neighbor Discovery	More than Two Routers Allowed in Subnet
Point-to-point	No	10	Yes	No
Broadcast	Yes	10	Yes	Yes
Nonbroadcast	Yes	30	No	Yes
Point-to-multipoint	No	30	Yes	Yes
Point-to-multipoint nonbroadcast	No	30	No	Yes
Loopback	No	—	—	No



OSPF Network Types

- **Point-to-point:** Routers use **multicast** to dynamically discover neighbors. There is no DR/BDR election. It is a default OSPF network type for serial links and point-to-point Frame Relay subinterfaces.
- **Broadcast:** **Multicast** is used to dynamically discover neighbors. The DR and BDR are elected. It is a default OSPF network type for Ethernet links.
- **Nonbroadcast:** Used on networks that interconnect more than two routers but without broadcast capability. **Neighbors must be statically configured**, followed by DR/BDR election. This network type is the default for all physical interfaces and multipoint subinterfaces using Frame Relay encapsulation.



OSPF Network Types

- **Point-to-multipoint:** OSPF treats this network type as a **logical collection of point-to-point links** even though all interfaces belong to the **common IP subnet**. Every interface IP address will appear in the routing table of the neighbors as a host /32 route. Neighbors are discovered dynamically using **multicast**. No DR/BDR election occurs.
- **Point-to-multipoint nonbroadcast:** Cisco extension that has the same characteristics as point-to-multipoint type except for the fact that **neighbors are not discovered** dynamically. Neighbors must be **statically defined**, and unicast is used for communication.
- **Loopback:** Default network type on loopback interfaces.



Configuring Passive Interfaces

```
Router(config)# router ospf 1
Router(config-if)# passive-interface default
Router(config-if)# no passive-interface serial 1/0
```

Building the Link- State Database



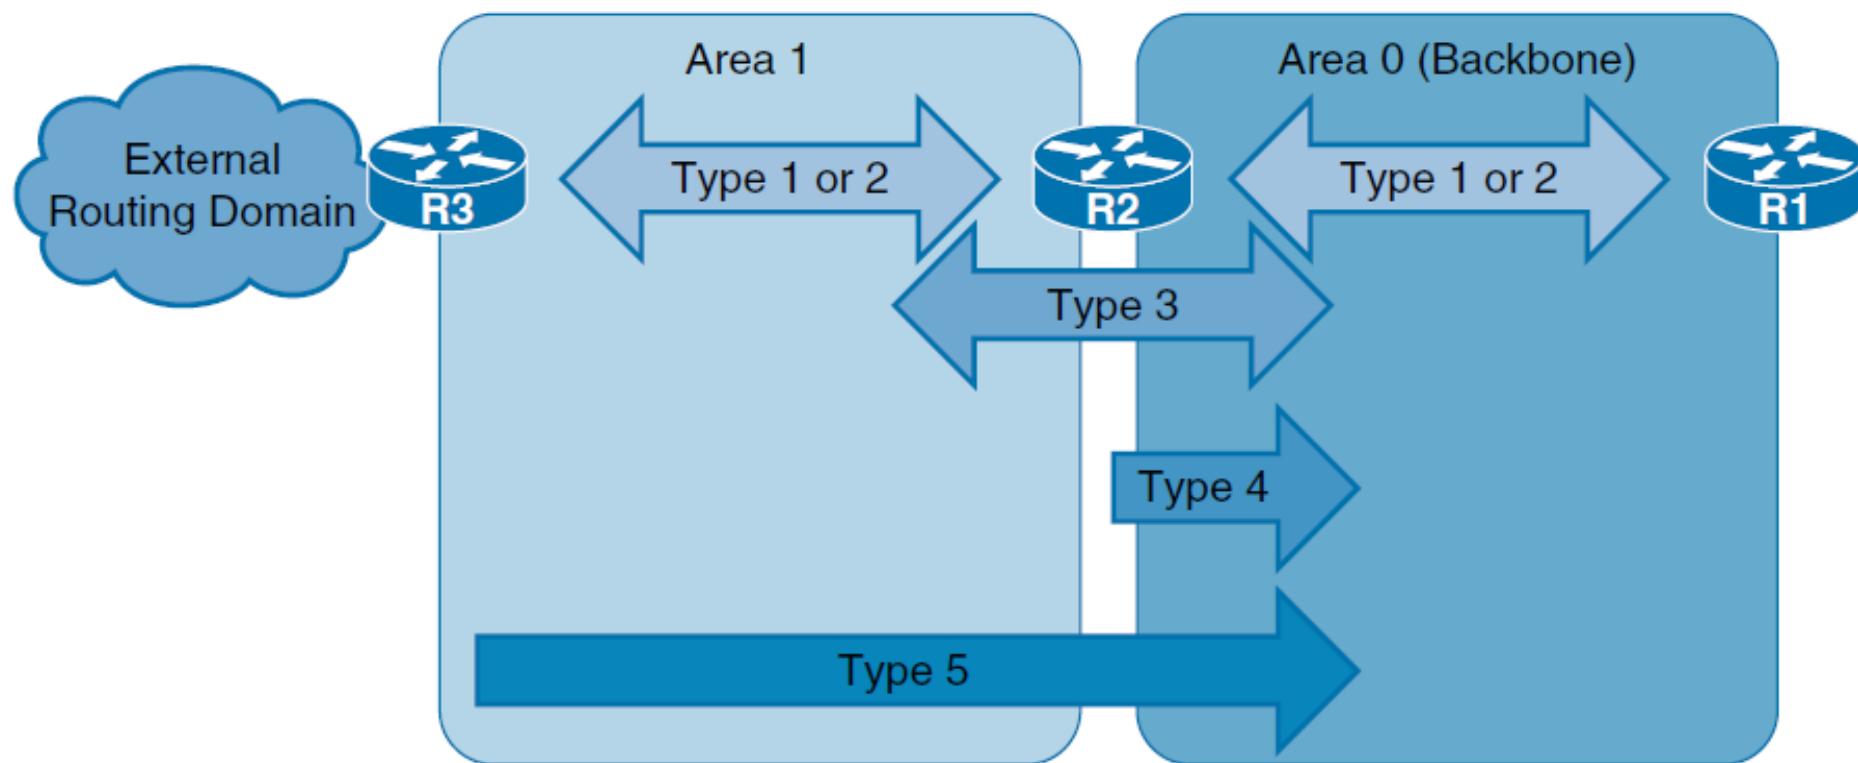


Building the Link-State Database

- List and describe different LSA types
- Describe how OSPF LSAs are also reflooded at periodic intervals
- Describe the exchange of information in a network without a designated router
- Describe the exchange of information in a network with a designated router
- Explain when SPF algorithms occur
- Describe how the cost of intra-area routes is calculated
- Describe how the cost of interarea routes is calculated
- Describe rules selecting between intra-area and interarea routes



OSPF LSA Types





OSPF LSA Types

- **Type 1, Router LSA:** Every router generates **router link advertisements** for each area to which it belongs. Router link advertisements describe the state of the router links to the area and are flooded only within that particular area. The **link-state ID** of the type 1 LSA is the **originating router ID**.
- **Type 2, Network LSA:** DRs generate **network link advertisements** for multi-access networks. Network link advertisements **describe the set of routers** that are attached to a particular multiaccess network. Network link advertisements are flooded in the area that contains the network. The **link-state ID** of the type 2 LSA is the **IP interface address of the DR**.
- **Type 3, Summary LSA:** An **ABR** takes the information that it learned in one area and **describes and summarizes** it for another area in the summary link advertisement. This **summarization is not on by default**. The **link-state ID** of the type 3 LSA is the destination network number.
- **Type 4, ASBR Summary LSA:** The ASBR summary link advertisement informs the rest of the OSPF domain **how to get to the ASBR**. The **link-state ID** includes the router ID of the described ASBR.
- **Type 5, Autonomous System LSA:** Autonomous system external link advertisements, which are generated by ASBRs, **describe routes to destinations that are external to the autonomous system**. They get flooded everywhere, except into special areas. The **link-state ID** of the type 5 LSA is the external network number.

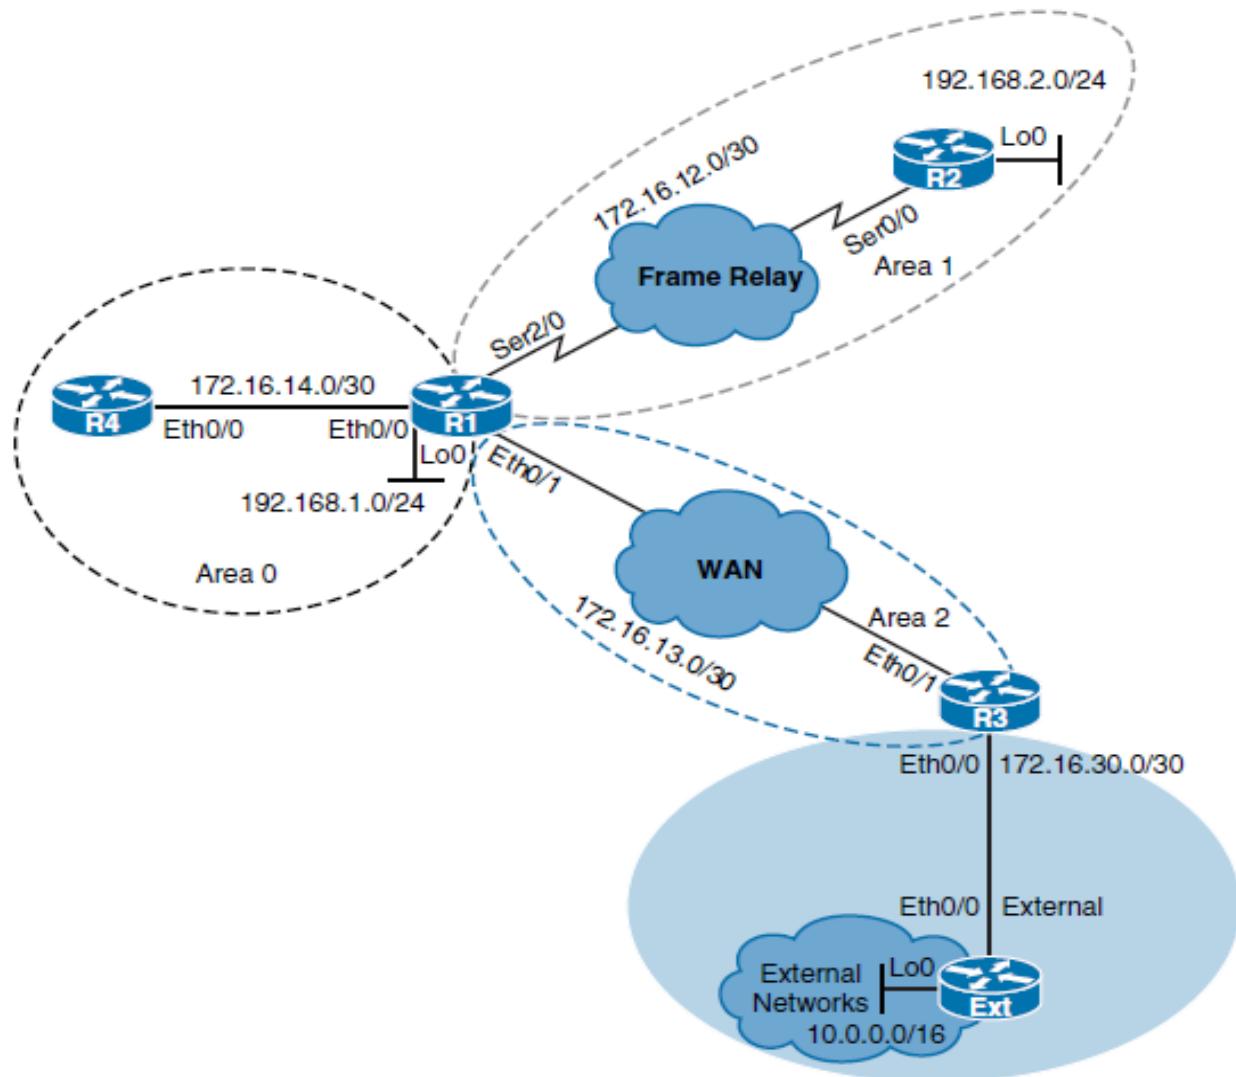


Other LSA Types

- **Type 6:** Specialized LSAs that are used in multicast OSPF applications
- **Type 7:** Used in special area type NSSA for external routes
- **Type 8, 9:** Used in OSPFv3 for link-local addresses and intra-area prefix
- **Type 10, 11:** Generic LSAs, also called *opaque* , which allow future extensions of OSPF



Examining the OSPF Link-State Database





Routing Table

```
R4# show ip route

Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, 1 - LISP
      + - replicated route, # - next hop override

Gateway of last resort is not set

      10.0.0.0/16 is subnetted, 1 subnets
O E2      10.0.0.0 [110/20] via 172.16.14.1, 00:46:48, Ethernet0/0
      172.16.0.0/16 is variably subnetted, 4 subnets, 2 masks
O IA      172.16.12.0/30 [110/74] via 172.16.14.1, 03:19:12, Ethernet0/0
O IA      172.16.13.0/30 [110/20] via 172.16.14.1, 03:19:12, Ethernet0/0
C        172.16.14.0/30 is directly connected, Ethernet0/0
L        172.16.14.2/32 is directly connected, Ethernet0/0
O        192.168.1.0/24 [110/11] via 172.16.14.1, 00:36:19, Ethernet0/0
O IA     192.168.2.0/24 [110/75] via 172.16.14.1, 00:47:59, Ethernet0/0
```



OSPF LSDB

```
R4# show ip ospf database
```

OSPF Router with ID (4.4.4.4) (Process ID 1)

LSA1

Router Link States (Area 0)

Link ID	ADV Router	Age	Seq#	Checksum	Link count
1.1.1.1	1.1.1.1	291	0x8000000B	0x00966C	2
4.4.4.4	4.4.4.4	1993	0x80000007	0x001C4E	1

LSA2

Net Link States (Area 0)

Link ID	ADV Router	Age	Seq#	Checksum
172.16.14.2	4.4.4.4	1993	0x80000006	0x0091B5

LSA3

Summary Net Link States (Area 0)

Link ID	ADV Router	Age	Seq#	Checksum
172.16.12.0	1.1.1.1	291	0x80000007	0x00C567
172.16.13.0	1.1.1.1	291	0x80000007	0x009CC5
192.168.2.0	1.1.1.1	1021	0x80000002	0x002FED

LSA4

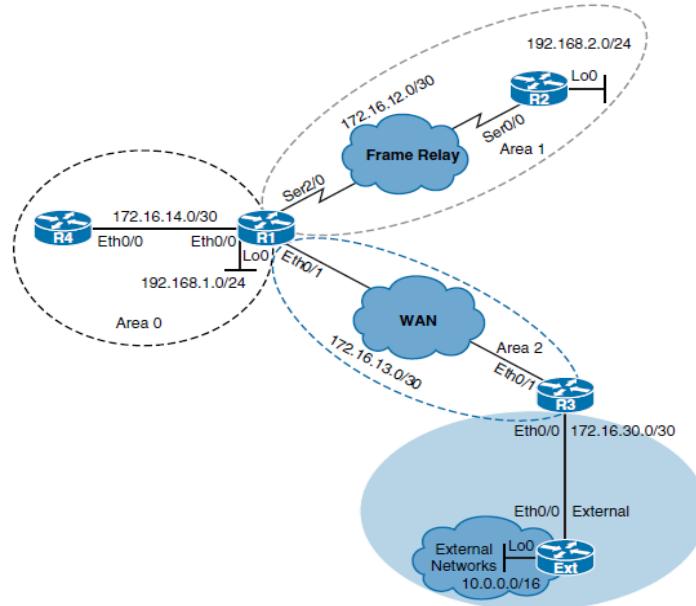
Summary ASB Link States (Area 0)

Link ID	ADV Router	Age	Seq#	Checksum
3.3.3.3	1.1.1.1	1031	0x80000002	0x0035EB

LSA5

Type-5 AS External Link States

Link ID	ADV Router	Age	Seq#	Checksum	Tag
10.0.0.0	3.3.3.3	977	0x80000002	0x000980	0





Type 1 LSA Details

```
R4# show ip ospf database router

    OSPF Router with ID (4.4.4.4) (Process ID 1)

        Router Link States (Area 0)

    Routing Bit Set on this LSA in topology Base with MTID 0

    LS age: 321
    Options: (No TOS-capability, DC)
    LS Type: Router Links
    Link State ID: 1.1.1.1
    Advertising Router: 1.1.1.1
    LS Seq Number: 8000000B
    Checksum: 0x966C
    Length: 48
    Area Border Router
    Number of Links: 2

    Link connected to: a Stub Network
    (Link ID) Network/subnet number: 192.168.1.0
```



Locally Generated Type 1 LSAs

```
R4# show ip ospf database router self-originate

    OSPF Router with ID (4.4.4.4) (Process ID 1)

        Router Link States (Area 0)

LS age: 23
Options: (No TOS-capability, DC)
LS Type: Router Links
Link State ID: 4.4.4.4
Advertising Router: 4.4.4.4
LS Seq Number: 80000008
Checksum: 0x1A4F
Length: 36
Number of Links: 1

Link connected to: a Transit Network
(Link ID) Designated Router address: 172.16.14.2
(Link Data) Router Interface address: 172.16.14.2
Number of MTID metrics: 0
TOS 0 Metrics: 10
```



Type 1 LSA

- Type 1 LSAs are generated by **every router** and **flooded within the area**.
- They describe the state of the router links in that area.
- When generating a type 1 LSA, the router uses its own **router ID as the value of LSID**.



OSPF Type 2 Network LSA

```
R4# show ip ospf database
```

```
OSPF Router with ID (4.4.4.4) (Process ID 1)
```

```
Router Link States (Area 0)
```

Link ID	ADV Router	Age	Seq#	Checksum	Link count
1.1.1.1	1.1.1.1	486	0x8000000B	0x00966C	2
4.4.4.4	4.4.4.4	142	0x80000008	0x001A4F	1

```
Net Link States (Area 0)
```

Link ID	ADV Router	Age	Seq#	Checksum
172.16.14.2	4.4.4.4	142	0x80000007	0x008FB6

<Output omitted>



OSPF Type 2 Network LSA

```
R4# show ip ospf database network

    OSPF Router with ID (4.4.4.4) (Process ID 1)

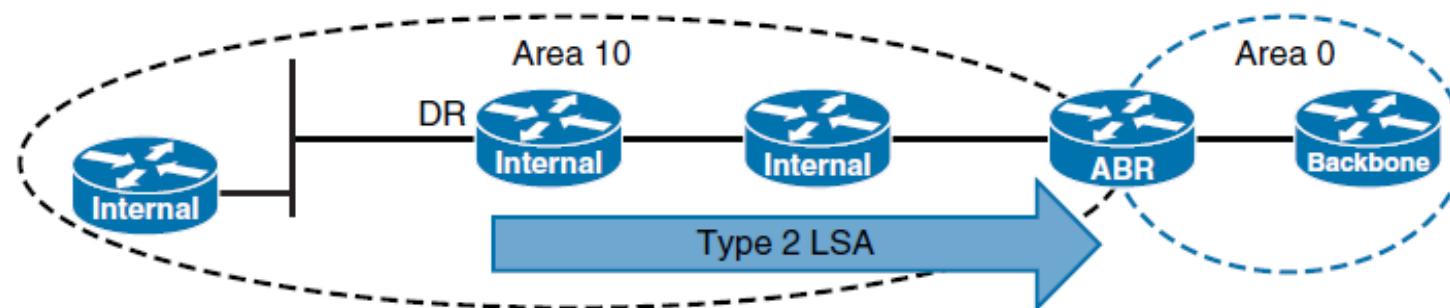
        Net Link States (Area 0)

Routing Bit Set on this LSA in topology Base with MTID 0
LS age: 170
Options: (No TOS-capability, DC)
LS Type: Network Links
Link State ID: 172.16.14.2 (address of Designated Router)
Advertising Router: 4.4.4.4
LS Seq Number: 80000007
Checksum: 0x8FB6
Length: 32
Network Mask: /30
Attached Router: 4.4.4.4
Attached Router: 1.1.1.1
```



Type 2 LSA

- The DR of the network is responsible for advertising the network LSA
- A type 2 **network LSA lists each of the attached routers** that make up the transit network, including the DR itself, and the subnet mask that is used on the link.
- The type 2 LSA then floods to all routers within the transit network area
- Type 2 LSAs never cross an area boundary.
- The LSID for a network LSA is the **IP interface address of the DR** that advertises it





OSPF Type 3 Summary LSA

```
R4# show ip ospf database
```

```
OSPF Router with ID (4.4.4.4) (Process ID 1)
```

```
Router Link States (Area 0)
```

Link ID	ADV Router	Age	Seq#	Checksum	Link count
1.1.1.1	1.1.1.1	583	0x8000000B	0x00966C	2
4.4.4.4	4.4.4.4	238	0x80000008	0x001A4F	1

```
Net Link States (Area 0)
```

Link ID	ADV Router	Age	Seq#	Checksum
172.16.14.2	4.4.4.4	238	0x80000007	0x008FB6

```
Summary Net Link States (Area 0)
```

Link ID	ADV Router	Age	Seq#	Checksum
172.16.12.0	1.1.1.1	583	0x80000007	0x00C567
172.16.13.0	1.1.1.1	583	0x80000007	0x009CC5
192.168.2.0	1.1.1.1	1322	0x80000002	0x002E5D

```
<Output omitted>
```



OSPF Type 3 Summary LSA

```
R4# show ip ospf database summary
```

```
OSPF Router with ID (4.4.4.4) (Process ID 1)
```

```
Summary Net Link States (Area 0)
```

Routing Bit Set on this LSA in topology Base with MTID 0

LS age: 608

Options: (No TOS-capability, DC, Upward)

LS Type: Summary Links (Network)

Link State ID: 172.16.12.0 (summary Network Number)

Advertising Router: 1.1.1.1

LS Seq Number: 80000007

Checksum: 0xC567

Length: 28

Network Mask: /30

MTID: 0 Metric: 64

Routing Bit Set on this LSA in topology Base with MTID 0

LS age: 608

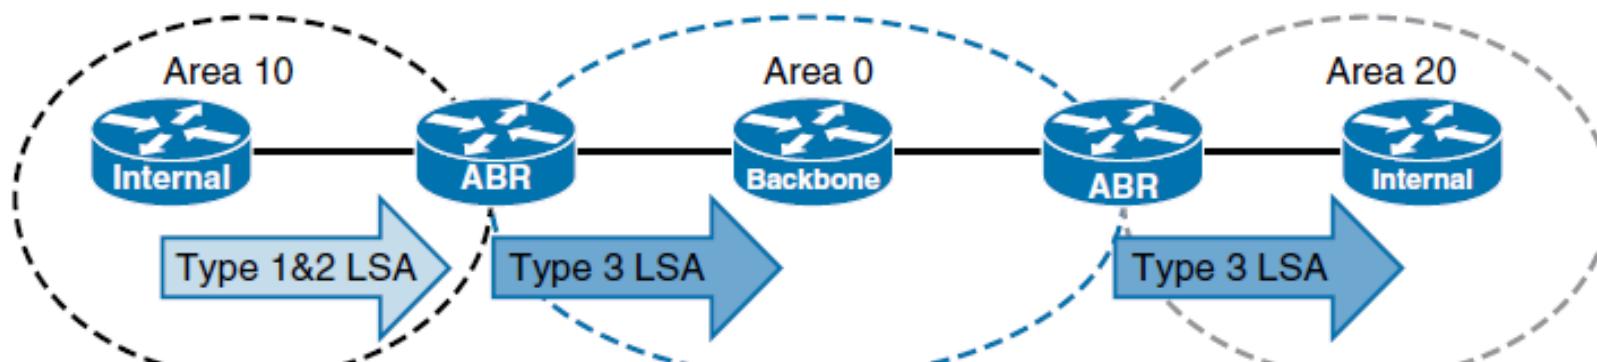
Options: (No TOS-capability, DC, Upward)

LS Type: Summary Links (Network)



Type 3 LSA

- The ABRs **generate type 3 summary LSAs** to describe any networks that are owned by an area to the rest of the areas in the OSPF autonomous system
- Summary LSAs are flooded throughout a single area only, but are regenerated by ABRs to flood into other areas.
- By default, OSPF does **not automatically summarize** groups of contiguous subnets.
- As a best practice, you can use **manual route summarization** on ABRs to limit the amount of information that is exchanged between the areas.





OSPF Type 4 ASBR Summary LSA

```
R4# show ip ospf database asbr-summary
```

```
OSPF Router with ID (4.4.4.4) (Process ID 1)
```

```
Summary ASB Link States (Area 0)
```

Routing Bit Set on this LSA in topology Base with MTID 0

LS age: 1420

Options: (No TOS-capability, DC, Upward)

LS Type: Summary Links (AS Boundary Router)

```
Link State ID: 3.3.3.3 (AS Boundary Router address)
```

```
Advertising Router: 1.1.1.1
```

LS Seq Number: 80000002

Checksum: 0x35EB

Length: 28

Network Mask: /0

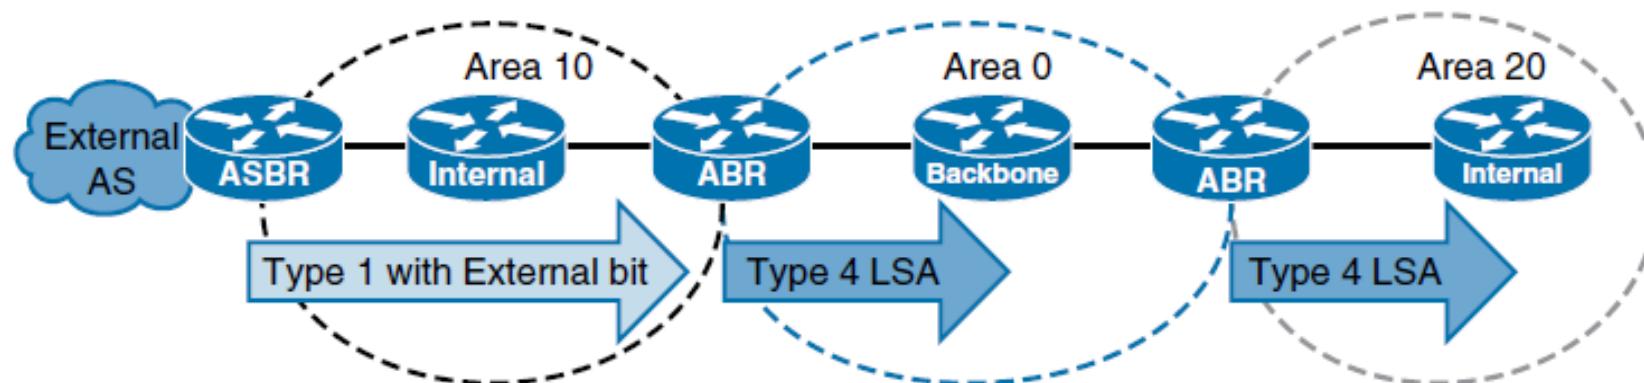
MTID: 0

Metric: 10



Type 4 LSA

- A type 4 LSA **identifies the ASBR** and provides a route to the ASBR.
- The link-state ID is set to the **ASBR router ID**.
- All traffic that is destined to an external autonomous system requires routing table knowledge of the ASBR that originated the external routes.





OSPF Type 5 External LSA

```
R4# show ip ospf database external

    OSPF Router with ID (4.4.4.4) (Process ID 1)

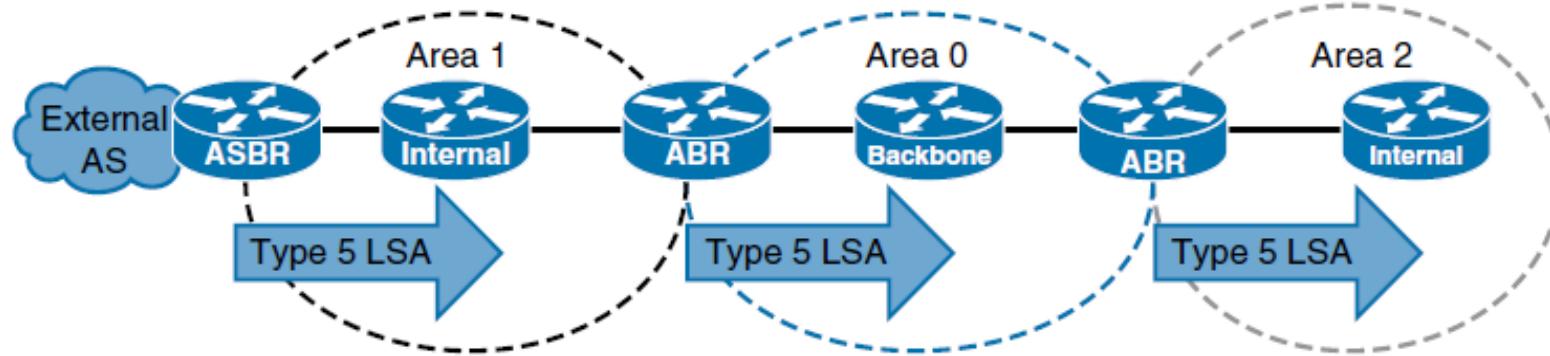
        Type-5 AS External Link States

Routing Bit Set on this LSA in topology Base with MTID 0
LS age: 1434
Options: (No TOS-capability, DC, Upward)
LS Type: AS External Link
Link State ID: 10.0.0.0 (External Network Number )
Advertising Router: 3.3.3.3
LS Seq Number: 80000002
Checksum: 0x980
Length: 36
Network Mask: /16
    Metric Type: 2 (Larger than any link state path)
    MTID: 0
    Metric: 20
    Forward Address: 0.0.0.0
External Route Tag: 0
```



Type 5 LSA

- The **LSID** is the external network number.
- Information described in the type 5 LSA combined with the information received in the type 4 LSA, which describes the ASBR capability of router R3. This way, R4 **learns** how to reach **the external networks**.





Periodic OSPF Database Changes

- Although OSPF does not refresh routing updates periodically, it does **reflood LSAs every 30 minutes**.
- Each LSA includes the **link-state age variable**, which counts the age of the LSA packet
- In a normally operating network, you will **not** see the **age** variable with values **higher than 1800** seconds.
- When an LSA reaches a max age of 60 minutes in the LSDB, it is removed from the LSDB, and the router will perform a new SPF calculation
- The router floods the LSA to other routers, informing them to remove the LSA as well.
- Because this update is only used to refresh the LSDB, it is sometimes called a ***paranoid update***



Exchanging and Synchronizing LSDBs

Neighbor state

Down



10.0.0.1

Link R1-R2 comes up

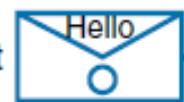
Neighbor state

192.168.1.1



Down

Init



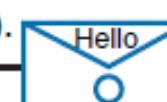
I am a router ID 10.0.0.1 and I see no one (multicast to 224.0.0.5).

Init

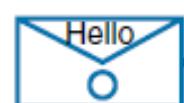
2-way



I am a router ID 10.0.0.2 and I see 10.0.0.1 (unicast to 10.0.0.1).



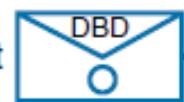
Init



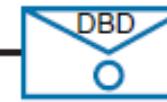
I am a router ID 10.0.0.1 and I see 10.0.0.2 (unicast to 10.0.0.2).

2-way

Exstart



I will start exchange because I have router ID 10.0.01.

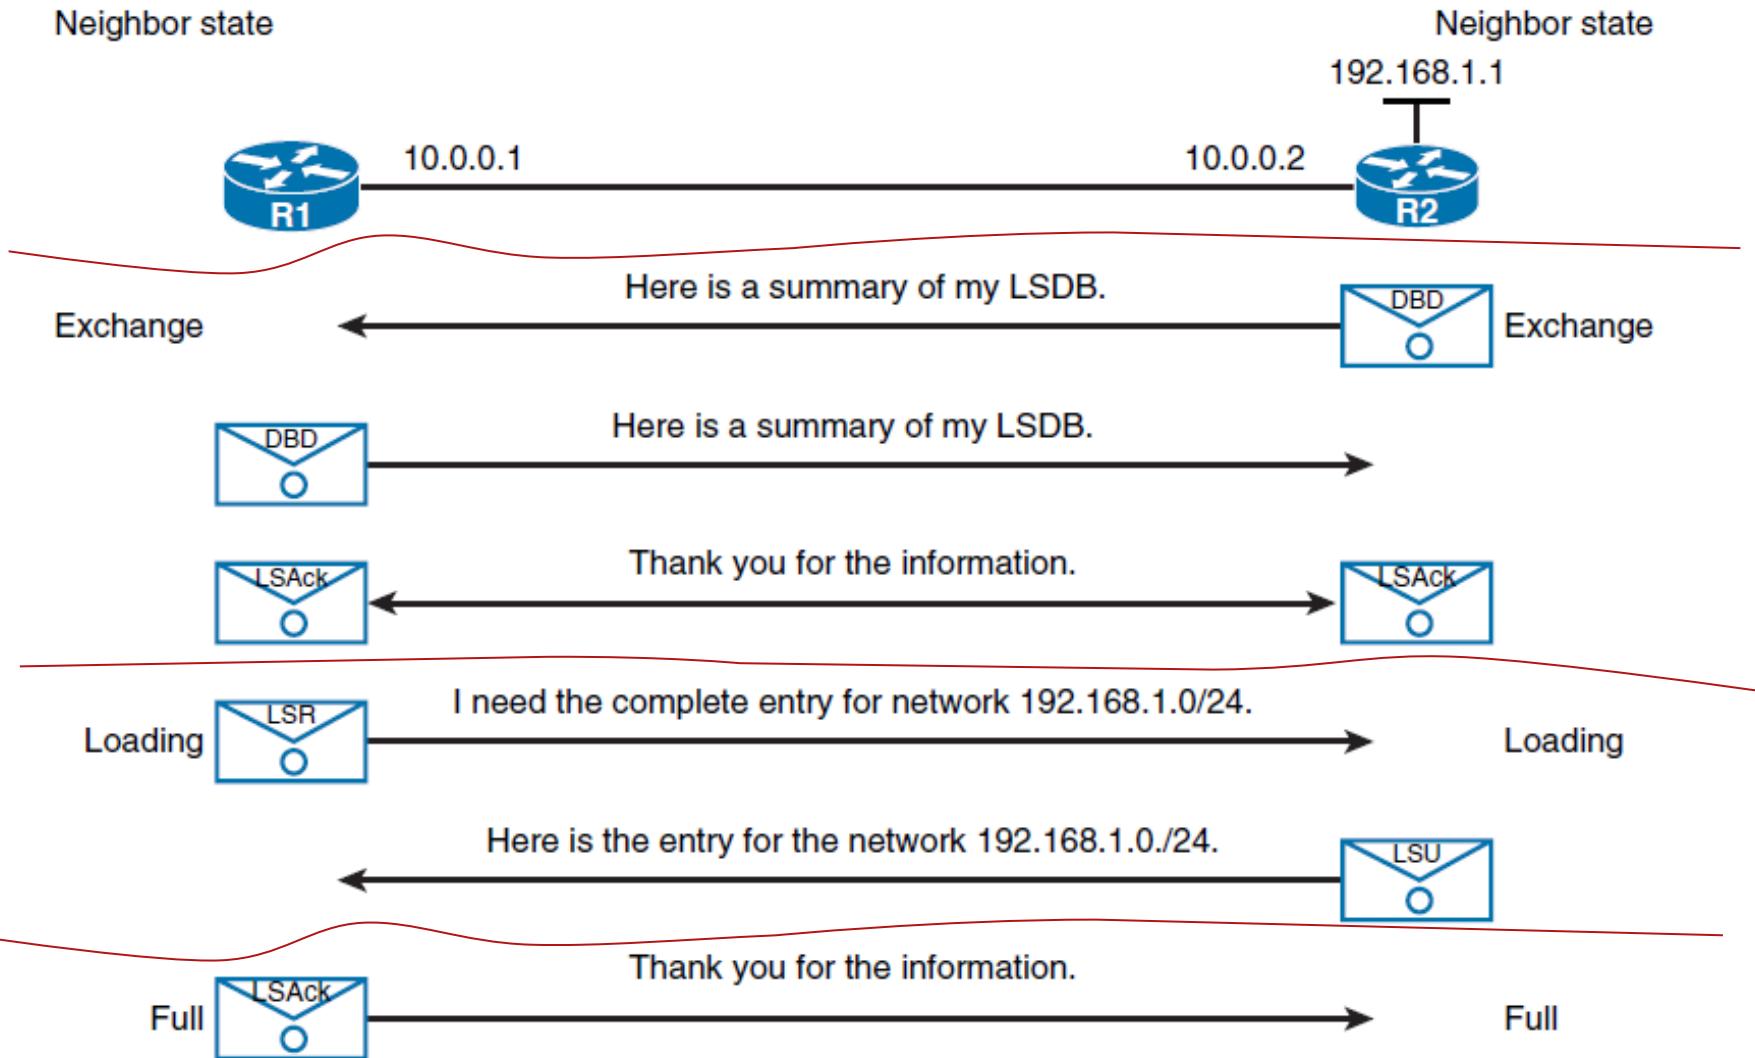


Exstart

No, I will start the exchange because I have a higher router ID.



Exchanging and Synchronizing LSDBs



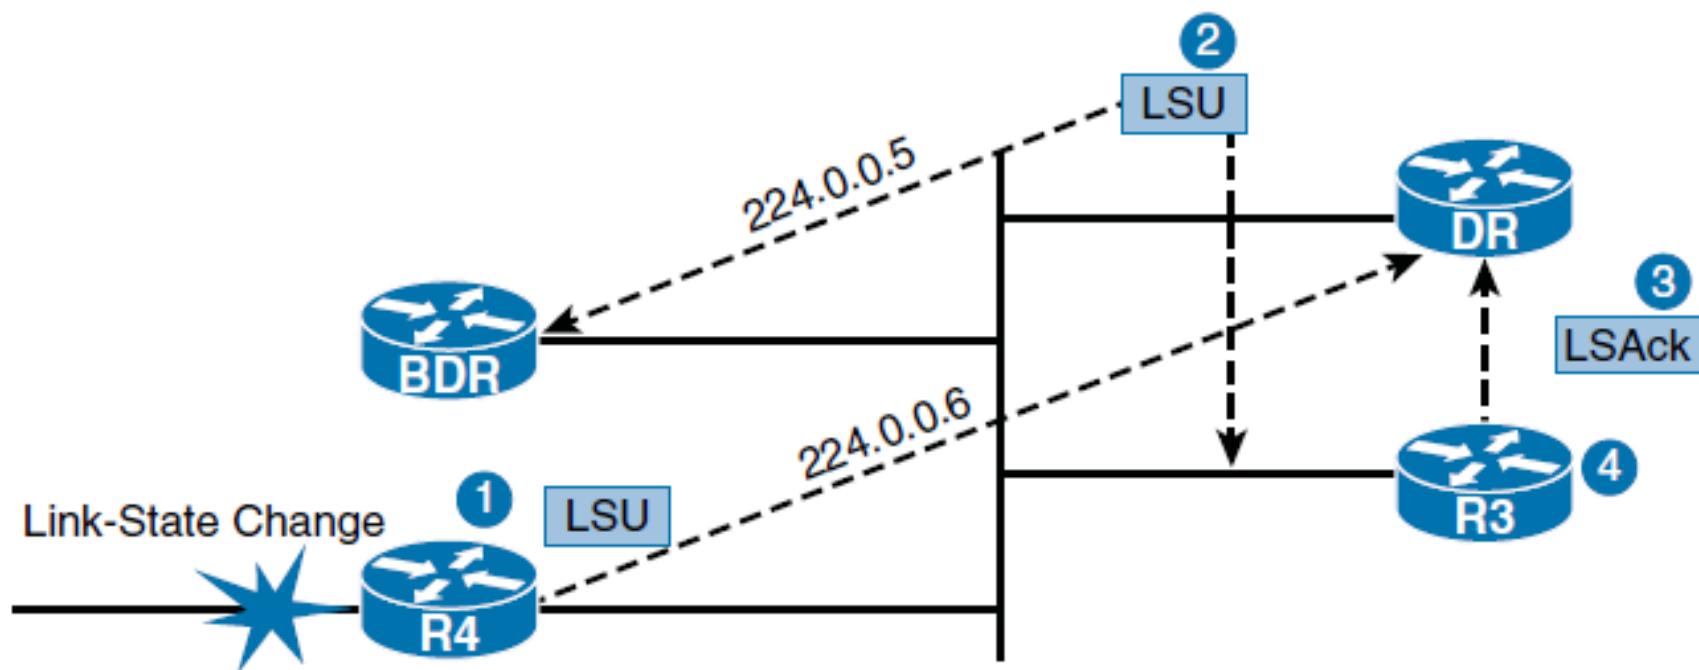


Synchronizing the LSDB on Multiaccess Networks

- When routers form a neighbor relationship on a multiaccess segment, the **DR and BDR election** takes place when routers are in the **2-Way state**.
- The router with a **highest OSPF priority**, or **highest router ID** in case of a tie, is elected as a DR. Similarly, the router with the second highest priority or router ID becomes the BDR.
- While the DR and BDR proceed in establishing the neighborship with all routers on the segment, other routers establish **full adjacency only with the DR and BDR**.
- The neighbor state of other neighbors stays in the **2-Way state**.
- Non-DR router exchange their databases only with the DR.
- The DR takes care to synchronize any new or changed LSAs with the rest of the routers on the segment.



Synchronizing the LSDB on Multiaccess Networks





Running the SPF Algorithm

- Every time there is a change in the network topology, OSPF needs to reevaluate its shortest path calculations. OSPF uses **SPF to determine best paths toward destinations.**
- The network topology that is described in the LSDB is used as an input for calculation.
- Network topology change can influence best path selection; therefore, routers must **rerun SPF each time there is an intra-area topology change.**
- Inter-area changes, which are described in **type 3 LSAs**, do **not trigger the SPF recalculation** because the input information for the best path calculation remains unchanged.
- The router determines the best paths for inter-area routes based on the calculation of the best path toward the ABR.



Verifying OSPF Frequency of the SPF Algorithm

```
R1# show ip ospf | begin Area
Area BACKBONE(0) (Inactive)
    Number of interfaces in this area is 1
    Area has no authentication
    SPF algorithm last executed 00:35:04:959 ago
    SPF algorithm executed 5 times
    Area ranges are
    Number of opaque link LSA 0. Checksum Sum 0x000000
```



OSPF Best Path Calculation

- Once LSDBs are synchronized among OSPF neighbors, each router needs to determine on its own the best paths over the network topology.
- SPF it **compares total costs** of specific paths against each other. The **paths with the lowest costs** are selected as **the best paths**.
- OSPF cost is computed automatically for each interface that is assigned into an OSPF process, using the following formula:

$$\text{Cost} = \text{Reference bandwidth} / \text{Interface bandwidth}$$

- The cost value is a **16-bit positive number between 1 and 65,535**, where a lower value is a more desirable metric.
- Reference bandwidth is set to 100 Mbps by default.
- On high-bandwidth links (100 Mbps and more), automatic cost assignment no longer works
- “On these links, OSPF costs must be set manually on each interface.”
- The OSPF cost is recomputed after every bandwidth change, and the **Dijkstra’s algorithm determines the best path by adding all link costs along a path**.



Examining the Interface Bandwidth and OSPF Cost

```
R1# show interface serial 2/0
```

Serial2/0 is up, line protocol is up

Hardware is M4T

Internet address is 172.16.12.1/30

MTU 1500 bytes, BW 1544 Kbit/sec, DLY 20000 usec,

reliability 255/255, txload 1/255, rxload 1/255

Encapsulation FRAME-RELAY, crc 16, loopback not set

<Output omitted>

```
R1# show ip ospf interface serial 2/0
```

Serial2/0 is up, line protocol is up

Internet Address 172.16.12.1/30, Area 1, Attached via Network Statement

Process ID 1, Router ID 1.1.1.1, Network Type NON_BROADCAST, Cost: 64

Topology-MTID	Cost	Disabled	Shutdown	Topology Name
---------------	------	----------	----------	---------------

0	64	no	no	Base
---	----	----	----	------

Transmit Delay is 1 sec, State BDR, Priority 1

Designated Router (ID) 2.2.2.2, Interface address 172.16.12.2

Backup Designated router (ID) 1.1.1.1, Interface address 172.16.12.1

Timer intervals configured, Hello 30, Dead 120, Wait 120, Retransmit 5

<Output omitted>



Default OSPF Costs

Link Type	Default Cost
T1 (1.544-Mbps serial link)	64
Ethernet	10
Fast Ethernet	1
Gigabit Ethernet	1
10-Gigabit Ethernet	1



Modifying the Reference Bandwidth

```
R1(config)# router ospf 1
R1(config-router)# auto-cost reference-bandwidth 10000
% OSPF: Reference bandwidth is changed.

    Please ensure reference bandwidth is consistent across all routers.
```

```
R1# show ip ospf interface serial 2/0
Serial2/0 is up, line protocol is up
    Internet Address 172.16.12.1/30, Area 1, Attached via Network Statement
    Process ID 1, Router ID 1.1.1.1, Network Type NON_BROADCAST, Cost: 6476
    Topology-MTID      Cost      Disabled      Shutdown      Topology Name
        0            6476        no           no           Base
<Output omitted>
```



Changing the Interface Bandwidth and OSPF Cost

```
R1(config)# interface serial 2/0
R1(config-if)# bandwidth 10000
```

```
R1# show interfaces serial 2/0
Serial2/0 is up, line protocol is up
  Hardware is M4T
  Internet address is 172.16.12.1/30
  MTU 1500 bytes, BW 10000 Kbit/sec, DLY 20000 usec,
<Output omitted>
R1# show ip ospf interface serial 2/0
Serial2/0 is up, line protocol is up
  Internet Address 172.16.12.1/30, Area 1, Attached via Network Statement
  Process ID 1, Router ID 1.1.1.1, Network Type NON_BROADCAST, Cost: 1000
  Topology-MTID      Cost      Disabled      Shutdown      Topology Name
                0        1000        no          no          Base
<Output omitted>
```

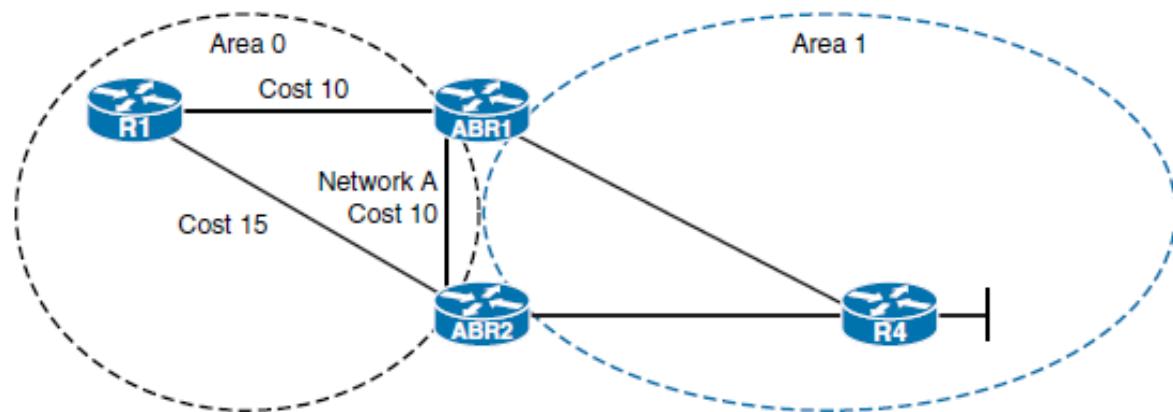
```
R1(config)# interface serial 2/0
R1(config-if)# ip ospf cost 500
```

```
R1# show ip ospf interface brief
Interface      PID      Area           IP Address/Mask      Cost      State Nbrs F/C
Lo0            1        0               192.168.1.1/24      1        P2P    0/0
Et0/0          1        0               172.16.14.1/30     1000     DR     1/1
Se2/0          1        1               172.16.12.1/30     500      BDR    1/1
Et0/1          1        2               172.16.13.1/30     1000     BDR    1/1
```



Calculating the Cost of Intra-Area Routes

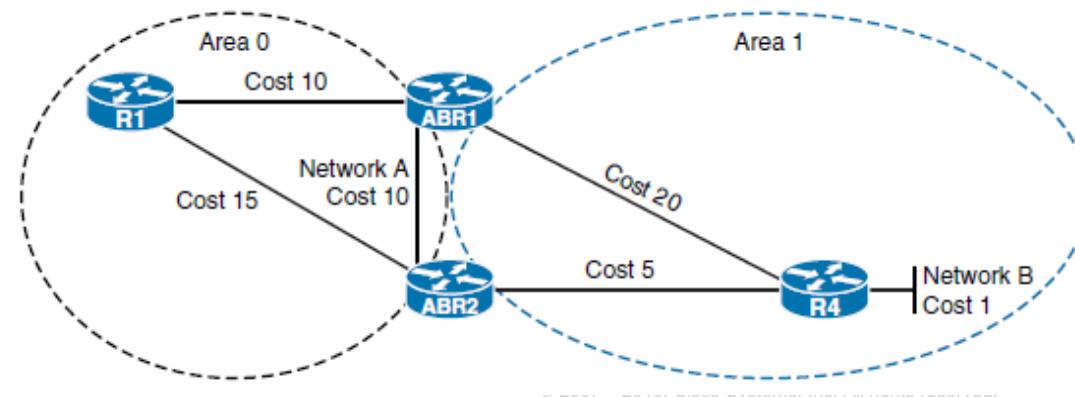
- To calculate the cost of intra-area routes, the router first analyzes OSPF database and identifies all subnets within its area.
- For each possible route, OSPF calculates the cost to reach the destination by summing up the individual interface costs.
- For each subnet, the route with the lowest total cost is selected as the best route.





Calculating the Cost of Inter-area Routes

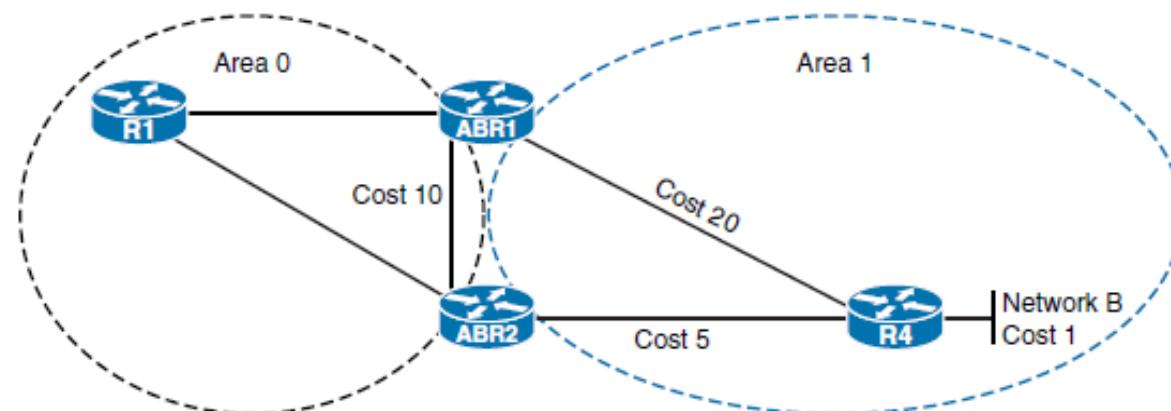
- The internal OSPF router within an area receives only summarized info about inter-area routes.
- As a result, the cost of an inter-area route cannot be calculated the same way as for the intra-area routes.
- When ABRs propagate information about the inter-area routes with type 3 LSAs, they include their lowest cost to reach a specific subnet in the advertisement.
- The **internal router adds its cost to reach a specific ABR** to the cost announced in a type 3 LSA.
- Then it selects the route with the lowest total cost as the best route.





Selecting Between Intra-Area and Inter-area Routes

- To eliminate the single point of failure on area borders, **at least two ABRs are used in most networks.**
- As a result, ABR can learn about a specific subnet from internal routers and also from the other ABR.
- ABR can learn an intra-area route and also an inter-area route for the same destination.
- Even though the **inter-area route could have lower cost** to the specific subnet, the **intra-area path is always the preferred choice.**



Optimizing OSPF Behavior





Optimizing OSPF Behavior

- Describe the properties of OSPF route summarization
- Describe benefits of route summarization in OSPF
- Configure summarization on ABR
- Configure summarization on ASBR
- Configure the cost of OSPF default route
- Describe how you can use default routes and stub routing to direct traffic toward the Internet
- Describe the NSSA areas
- Configure the default route using the **default-information originate** command



OSPF Route Summarization

- Route summarization is a key to scalability in OSPF.
- Route summarization helps solve two major problems:
 - Large routing tables
 - Frequent LSA flooding throughout the autonomous system
- Every time that a route disappears in one area, routers in other areas also get involved in shortest-path calculation.
- To reduce the size of the area database, you can configure summarization on an area boundary or autonomous system boundary.



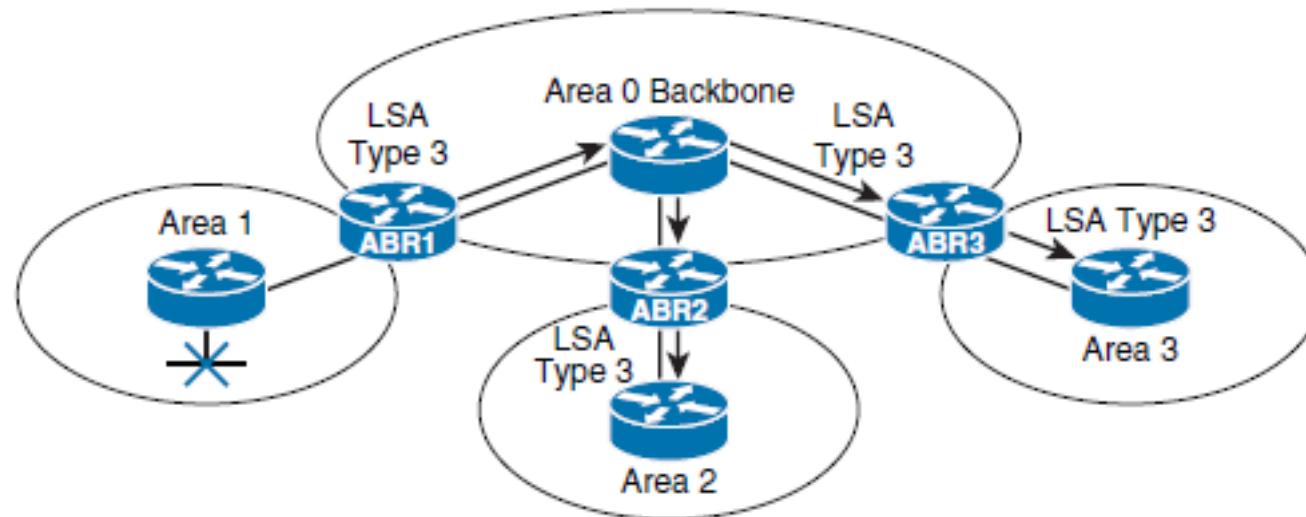
OSPF Route Summarization Process

- Normally, type 1 and type 2 LSAs are generated inside each area and translated into type 3 LSAs in other areas.
- With route summarization, the ABRs or ASBRs consolidate multiple routes into a single advertisement. ABRs summarize type 3 LSAs, and ASBRs summarize type 5 LSAs.
- Instead of advertising many specific prefixes, **advertise only one summary prefix**.
- Route summarization requires a good addressing plan—an assignment of subnets and addresses that is based on the OSPF area structure and lends itself to aggregation at the OSPF area borders.



Benefits of Route Summarization

- Route summarization directly affects the amount of bandwidth, CPU power, and memory resources that the OSPF routing process consumes.
- With route summarization, only the summarized routes are propagated into the backbone (area 0) increasing the stability of the network
- The routes being advertised in the type 3 LSAs are appropriately added to or deleted from the router's routing table, but an SPF calculation is not done.



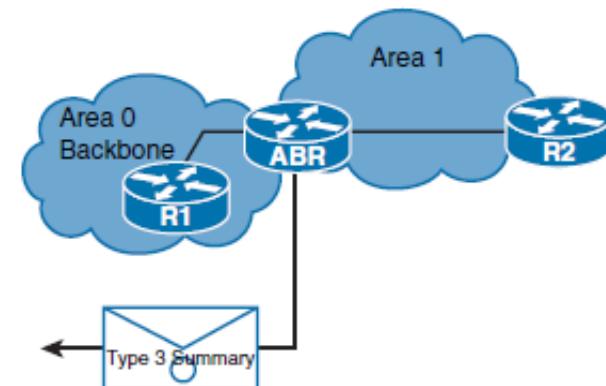


OSPF Route Summarization

- OSPF offers two methods of route summarization:
 - Summarization of internal routes performed on the ABRs
 - Summarization of external routes performed on the ASBRs



ABR's Summarization

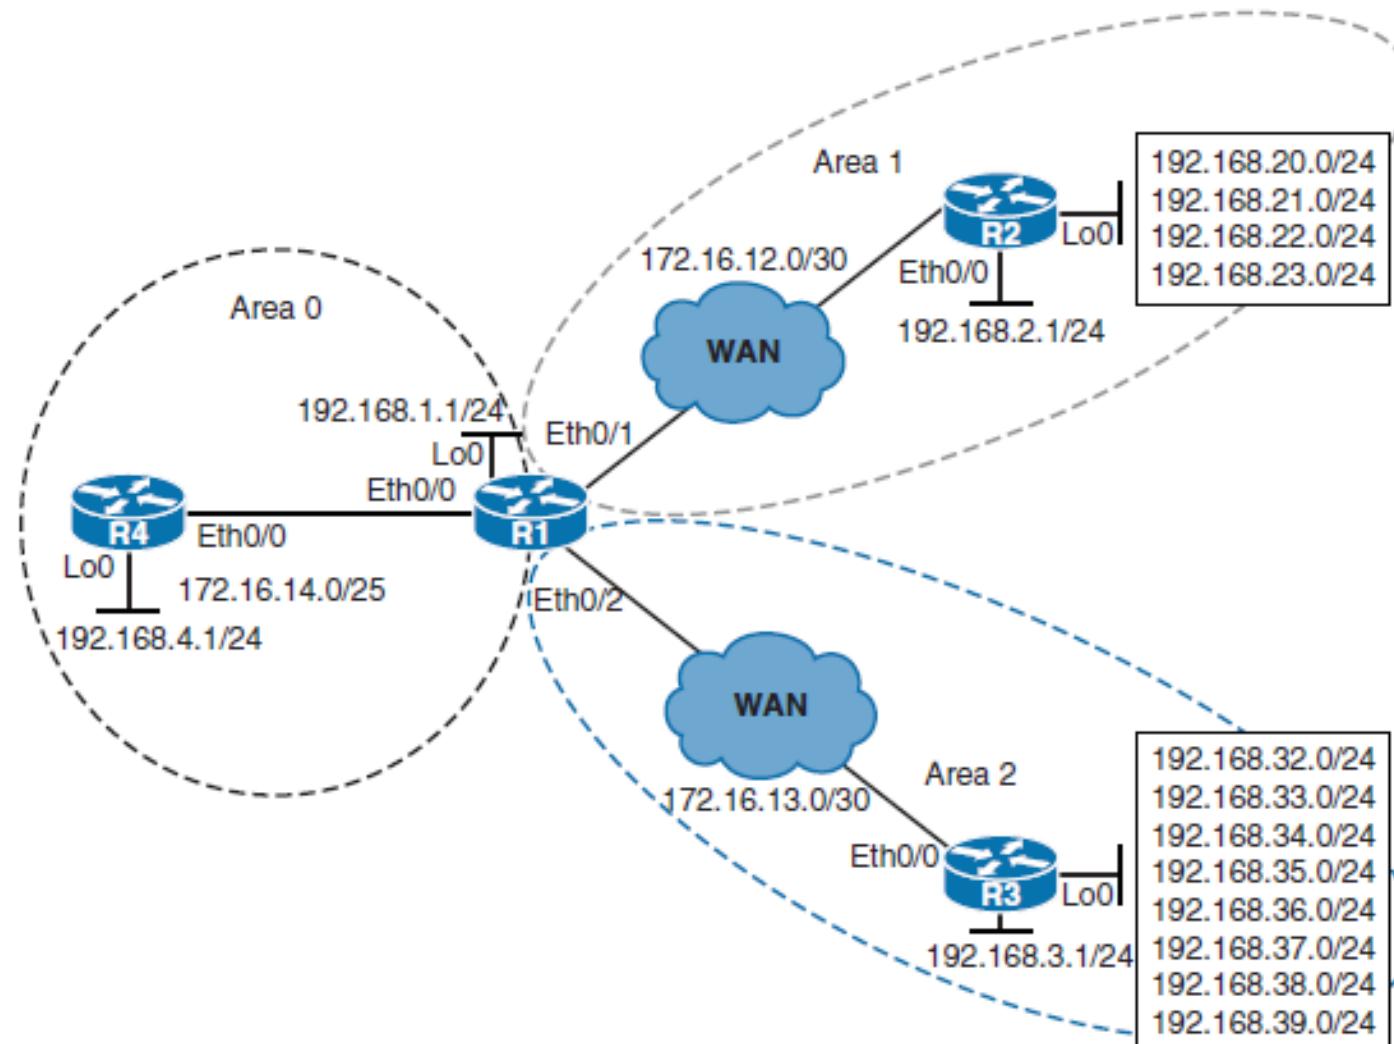


Parameter	Description
<i>area-id</i>	Identifier of the area about which routes are to be summarized. It can be specified as either a decimal value or as an IP address.
<i>ip-address</i>	IP address.
<i>mask</i>	IP address mask.
<i>advertise</i>	(Optional) Sets the address range status to advertise and generates a type 3 summary LSA.
<i>not-advertise</i>	(Optional) Sets the address range status to DoNotAdvertise. The Type 3 summary LSA is suppressed, and the component networks remain hidden from other networks.
<i>cost cost</i>	(Optional) Metric or cost for this summary route, which is used during OSPF SPF calculation to determine the shortest paths to the destination. The value can be 0 to 16,777,215.

- **`area area-id range ip-address mask [advertise | not-advertise] [cost cost]`**



Configuring OSPF Route Summarization





Configuring OSPF Route Summarization

```
R1# show ip route ospf
<Output omitted>

O 192.168.2.0/24 [110/11] via 172.16.12.2, 00:41:47, Ethernet0/1
O 192.168.3.0/24 [110/11] via 172.16.13.2, 00:40:01, Ethernet0/2
O 192.168.4.0/24 [110/11] via 172.16.14.2, 00:38:09, Ethernet0/0
O 192.168.20.0/24 [110/11] via 172.16.12.2, 00:41:37, Ethernet0/1
O 192.168.21.0/24 [110/11] via 172.16.12.2, 01:03:46, Ethernet0/1
O 192.168.22.0/24 [110/11] via 172.16.12.2, 01:03:36, Ethernet0/1
O 192.168.23.0/24 [110/11] via 172.16.12.2, 01:03:26, Ethernet0/1
O 192.168.32.0/24 [110/11] via 172.16.13.2, 00:40:14, Ethernet0/2
O 192.168.33.0/24 [110/11] via 172.16.13.2, 00:57:01, Ethernet0/2
O 192.168.34.0/24 [110/11] via 172.16.13.2, 00:01:16, Ethernet0/2
O 192.168.35.0/24 [110/11] via 172.16.13.2, 00:01:06, Ethernet0/2
O 192.168.36.0/24 [110/11] via 172.16.13.2, 00:00:56, Ethernet0/2
O 192.168.37.0/24 [110/11] via 172.16.13.2, 00:00:46, Ethernet0/2
O 192.168.38.0/24 [110/11] via 172.16.13.2, 00:00:32, Ethernet0/2
O 192.168.39.0/24 [110/11] via 172.16.13.2, 00:00:18, Ethernet0/2
```



Configuring OSPF Route Summarization

```
R4# show ip route ospf
<Output omitted>

      172.16.0.0/16 is variably subnetted, 4 subnets, 3 masks
O IA    172.16.12.0/30 [110/20] via 172.16.14.1, 01:17:30, Ethernet0/0
O IA    172.16.13.0/30 [110/20] via 172.16.14.1, 01:17:30, Ethernet0/0
O     192.168.1.0/24 [110/11] via 172.16.14.1, 01:17:30, Ethernet0/0
O IA   192.168.2.0/24 [110/21] via 172.16.14.1, 00:49:23, Ethernet0/0
O IA   192.168.3.0/24 [110/21] via 172.16.14.1, 00:47:37, Ethernet0/0
O IA   192.168.20.0/24 [110/21] via 172.16.14.1, 00:49:08, Ethernet0/0
O IA   192.168.21.0/24 [110/21] via 172.16.14.1, 01:11:23, Ethernet0/0
O IA   192.168.22.0/24 [110/21] via 172.16.14.1, 01:11:13, Ethernet0/0
O IA   192.168.23.0/24 [110/21] via 172.16.14.1, 01:11:03, Ethernet0/0
O IA   192.168.32.0/24 [110/21] via 172.16.14.1, 00:47:50, Ethernet0/0
O IA   192.168.33.0/24 [110/21] via 172.16.14.1, 01:04:37, Ethernet0/0
O IA   192.168.34.0/24 [110/21] via 172.16.14.1, 00:02:26, Ethernet0/0
O IA   192.168.35.0/24 [110/21] via 172.16.14.1, 00:02:16, Ethernet0/0
O IA   192.168.36.0/24 [110/21] via 172.16.14.1, 00:02:06, Ethernet0/0
O IA   192.168.37.0/24 [110/21] via 172.16.14.1, 00:01:56, Ethernet0/0
O IA   192.168.38.0/24 [110/21] via 172.16.14.1, 00:01:43, Ethernet0/0
O IA   192.168.39.0/24 [110/21] via 172.16.14.1, 00:01:28, Ethernet0/0
```



Configuring OSPF Route Summarization

```
R4# show ip ospf database
```

OSPF Router with ID (4.4.4.4) (Process ID 1)

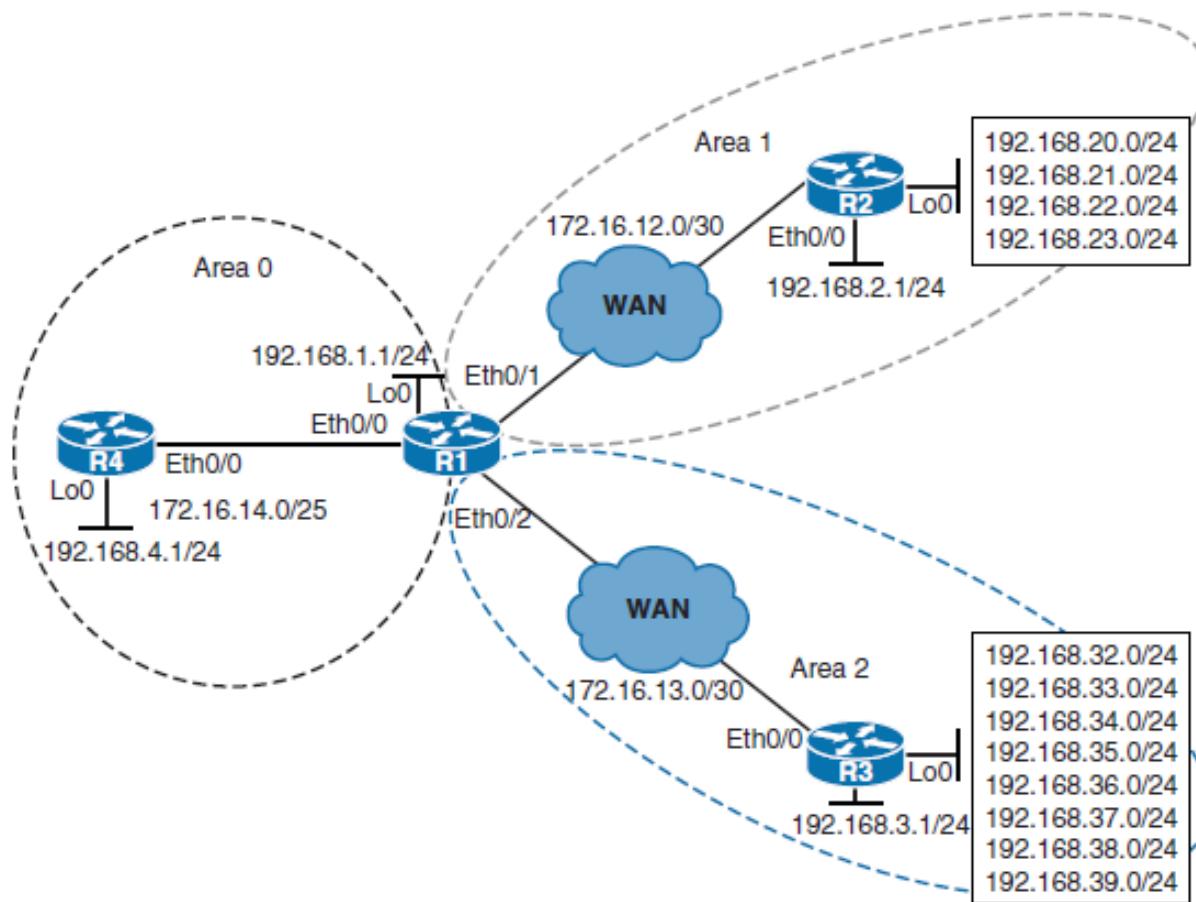
Summary Net Link States (Area 0)

Link ID	ADV Router	Age	Seq#	Checksum
172.16.12.0	1.1.1.1	553	0x80000008	0x00A5BC
172.16.13.0	1.1.1.1	553	0x80000008	0x009AC6
192.168.2.0	1.1.1.1	1541	0x80000006	0x0008B5
192.168.3.0	1.1.1.1	3607	0x80000007	0x008C3A
192.168.20.0	1.1.1.1	1541	0x8000000B	0x00376F
192.168.21.0	1.1.1.1	1800	0x80000004	0x003A72
192.168.22.0	1.1.1.1	1800	0x80000004	0x002F7C
192.168.23.0	1.1.1.1	1800	0x80000004	0x002486
192.168.32.0	1.1.1.1	3607	0x80000007	0x004C5D
192.168.33.0	1.1.1.1	3607	0x80000008	0x003F68
192.168.34.0	1.1.1.1	3607	0x80000002	0x00406C
192.168.35.0	1.1.1.1	3607	0x80000002	0x003576
192.168.36.0	1.1.1.1	3607	0x80000002	0x002A80
192.168.37.0	1.1.1.1	3607	0x80000002	0x001F8A
192.168.38.0	1.1.1.1	3607	0x80000002	0x001494
192.168.39.0	1.1.1.1	3607	0x80000002	0x00099E



Configuring Summarization on the ABR

```
R1(config)# router ospf 1  
R1(config-router)# area 1 range 192.168.20.0 255.255.252.0  
R1(config-router)# area 2 range 192.168.32.0 255.255.248.0
```





Configuring OSPF Route Summarization

```
R2# show ip route ospf
<Output omitted>

      172.16.0.0/16 is variably subnetted, 4 subnets, 3 masks
O IA    172.16.13.0/30 [110/20] via 172.16.12.1, 05:27:05, Ethernet0/0
O IA    172.16.14.0/25 [110/20] via 172.16.12.1, 05:07:35, Ethernet0/0
O IA  192.168.1.0/24 [110/11] via 172.16.12.1, 05:27:09, Ethernet0/0
O IA  192.168.3.0/24 [110/21] via 172.16.12.1, 01:24:16, Ethernet0/0
O IA  192.168.4.0/24 [110/21] via 172.16.12.1, 04:32:02, Ethernet0/0
O IA  192.168.32.0/21 [110/21] via 172.16.12.1, 00:57:42, Ethernet0/0
```

```
R3# show ip route ospf
<Output omitted>

      172.16.0.0/16 is variably subnetted, 4 subnets, 3 masks
O IA    172.16.12.0/30 [110/20] via 172.16.13.1, 05:25:50, Ethernet0/0
O IA    172.16.14.0/25 [110/20] via 172.16.13.1, 05:10:02, Ethernet0/0
O IA  192.168.1.0/24 [110/11] via 172.16.13.1, 05:25:50, Ethernet0/0
O IA  192.168.2.0/24 [110/21] via 172.16.13.1, 04:38:07, Ethernet0/0
O IA  192.168.4.0/24 [110/21] via 172.16.13.1, 04:34:29, Ethernet0/0
O IA  192.168.20.0/22 [110/21] via 172.16.13.1, 01:00:19, Ethernet0/0
```



Configuring OSPF Route Summarization

```
R4# show ip route ospf
<Output omitted>

    172.16.0.0/16 is variably subnetted, 4 subnets, 3 masks
O IA      172.16.12.0/30 [110/20] via 172.16.14.1, 05:16:24, Ethernet0/0
O IA      172.16.13.0/30 [110/20] via 172.16.14.1, 05:16:24, Ethernet0/0
O         192.168.1.0/24 [110/11] via 172.16.14.1, 05:16:24, Ethernet0/0
O IA      192.168.2.0/24 [110/21] via 172.16.14.1, 04:48:17, Ethernet0/0
O IA      192.168.3.0/24 [110/21] via 172.16.14.1, 01:36:53, Ethernet0/0
O IA      192.168.20.0/22 [110/21] via 172.16.14.1, 01:10:29, Ethernet0/0
O IA      192.168.32.0/21 [110/21] via 172.16.14.1, 01:10:19, Ethernet0/0
```



Configuring OSPF Route Summarization

```
R4# show ip ospf database  
<Output omitted>
```

Summary Net Link States (Area 0)

Link ID	ADV Router	Age	Seq#	Checksum
172.16.12.0	1.1.1.1	599	0x8000000B	0x009FBF
172.16.13.0	1.1.1.1	599	0x8000000B	0x0094C9
192.168.2.0	1.1.1.1	1610	0x80000009	0x0002B8
192.168.3.0	1.1.1.1	98	0x80000004	0x0001BD
192.168.20.0	1.1.1.1	599	0x8000000F	0x002085
192.168.32.0	1.1.1.1	98	0x80000005	0x009B0C



Configuring OSPF Route Summarization

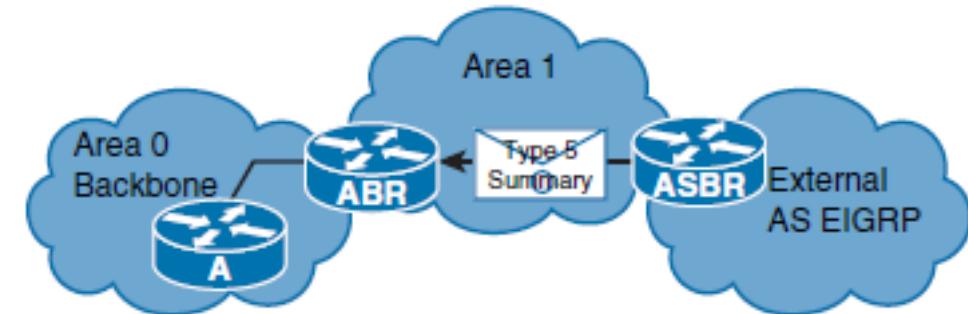
```
R1# show ip route ospf
```

```
<Output omitted>
```

```
O 192.168.2.0/24 [110/11] via 172.16.12.2, 01:18:25, Ethernet0/1
O 192.168.3.0/24 [110/11] via 172.16.13.2, 01:18:25, Ethernet0/2
O 192.168.4.0/24 [110/11] via 172.16.14.2, 01:18:25, Ethernet0/0
O 192.168.20.0/22 is a summary, 01:18:25, Null0
O 192.168.20.0/24 [110/11] via 172.16.12.2, 01:18:25, Ethernet0/1
O 192.168.21.0/24 [110/11] via 172.16.12.2, 01:18:25, Ethernet0/1
O 192.168.22.0/24 [110/11] via 172.16.12.2, 01:18:25, Ethernet0/1
O 192.168.23.0/24 [110/11] via 172.16.12.2, 01:18:25, Ethernet0/1
O 192.168.32.0/21 is a summary, 01:18:25, Null0
O 192.168.32.0/24 [110/11] via 172.16.13.2, 01:18:25, Ethernet0/2
O 192.168.33.0/24 [110/11] via 172.16.13.2, 01:18:25, Ethernet0/2
O 192.168.34.0/24 [110/11] via 172.16.13.2, 01:18:25, Ethernet0/2
O 192.168.35.0/24 [110/11] via 172.16.13.2, 01:18:25, Ethernet0/2
O 192.168.36.0/24 [110/11] via 172.16.13.2, 01:18:25, Ethernet0/2
O 192.168.37.0/24 [110/11] via 172.16.13.2, 01:18:25, Ethernet0/2
O 192.168.38.0/24 [110/11] via 172.16.13.2, 01:18:25, Ethernet0/2
O 192.168.39.0/24 [110/11] via 172.16.13.2, 01:18:25, Ethernet0/2
```



Summarization on ASBR's



Parameter	Description
<i>ip-address</i>	Summary address designated for a range of addresses.
<i>mask</i>	IP subnet mask used for the summary route.
<i>prefix</i>	IP route prefix for the destination.
<i>mask</i>	IP subnet mask used for the summary route.
<i>not-advertise</i>	(Optional) Suppress routes that match the specified prefix/mask pair. This keyword applies to OSPF only.
<i>tag tag</i>	(Optional) Tag value that can be used as a “match” value for controlling redistribution via route maps. This keyword applies to OSPF only.

- **summary-address** {{ *ip-address mask* } | { *prefix mask* }} [**not-advertise**] [**tag tag**]

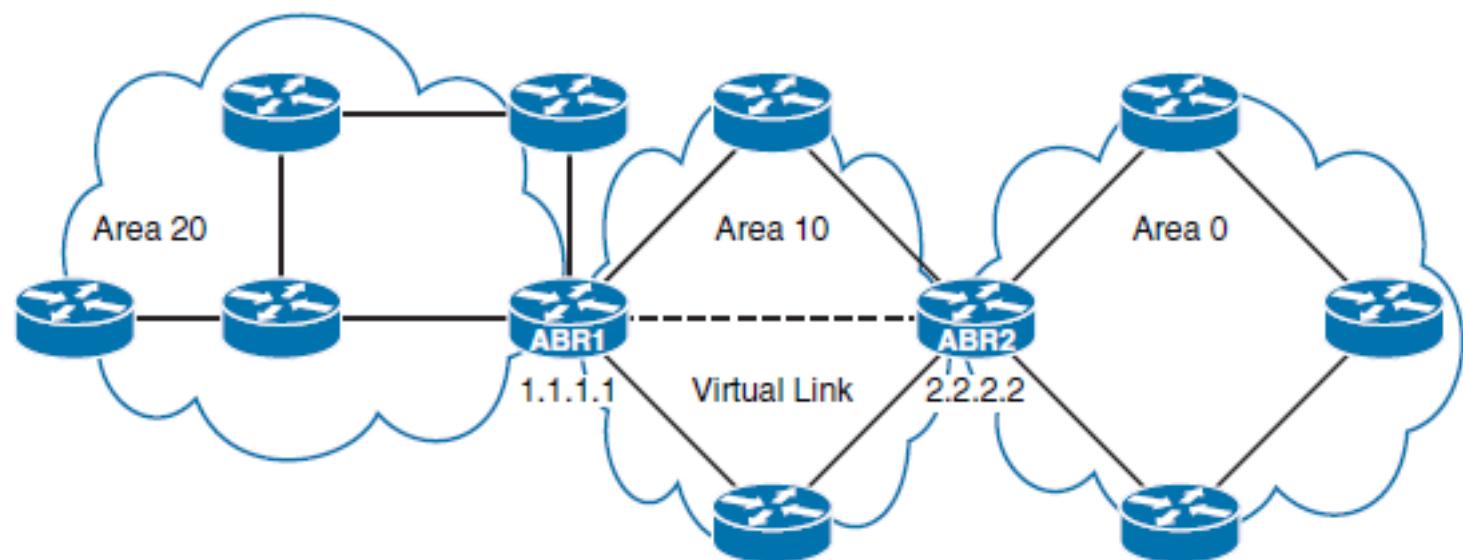
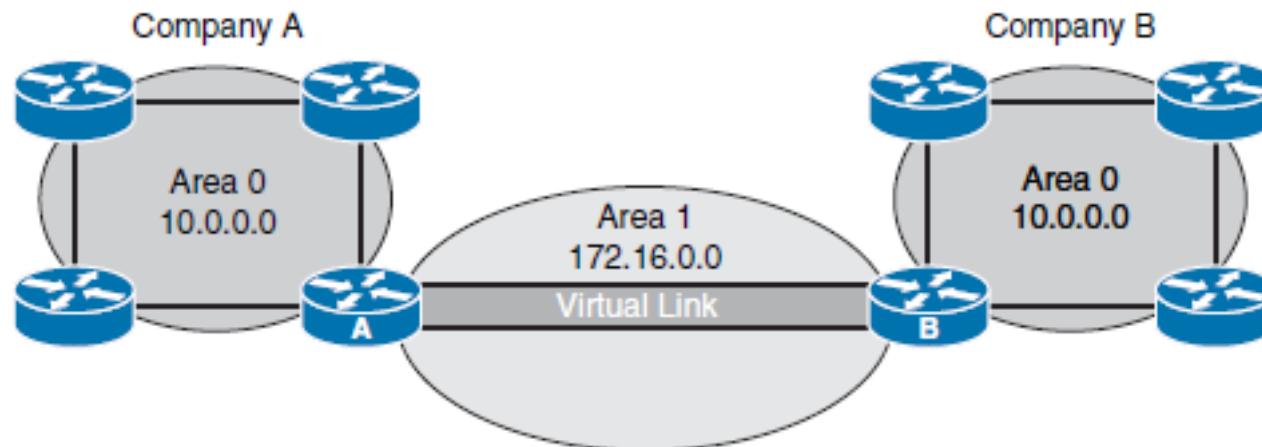


OSPF Virtual Links

- OSPF's two-tiered area hierarchy requires that if more than one area is configured, one of the areas **must be area 0**, the backbone area.
- All other areas must be directly connected to area 0, and area 0 must be contiguous.
- OSPF expects all nonbackbone areas to inject routes into the backbone, so that the routes can be distributed to other areas.
- A virtual link is a link that allows **discontiguous area 0s** to be connected, or a **disconnected area** to be connected to area 0, via a transit area.



OSPF Virtual Links





Virtual Links Issues

- The OSPF virtual link feature should be used **only in very specific cases**, for temporary connections or for backup after a failure.
- Virtual links should not be used as a primary backbone design feature.
- The virtual link relies on the stability of the underlying intra-area routing.
- Virtual links **cannot go through more than one area**, nor through stub areas.
- If a virtual link needs to be attached to the backbone across two nonbackbone areas, two virtual links are required, one per area.

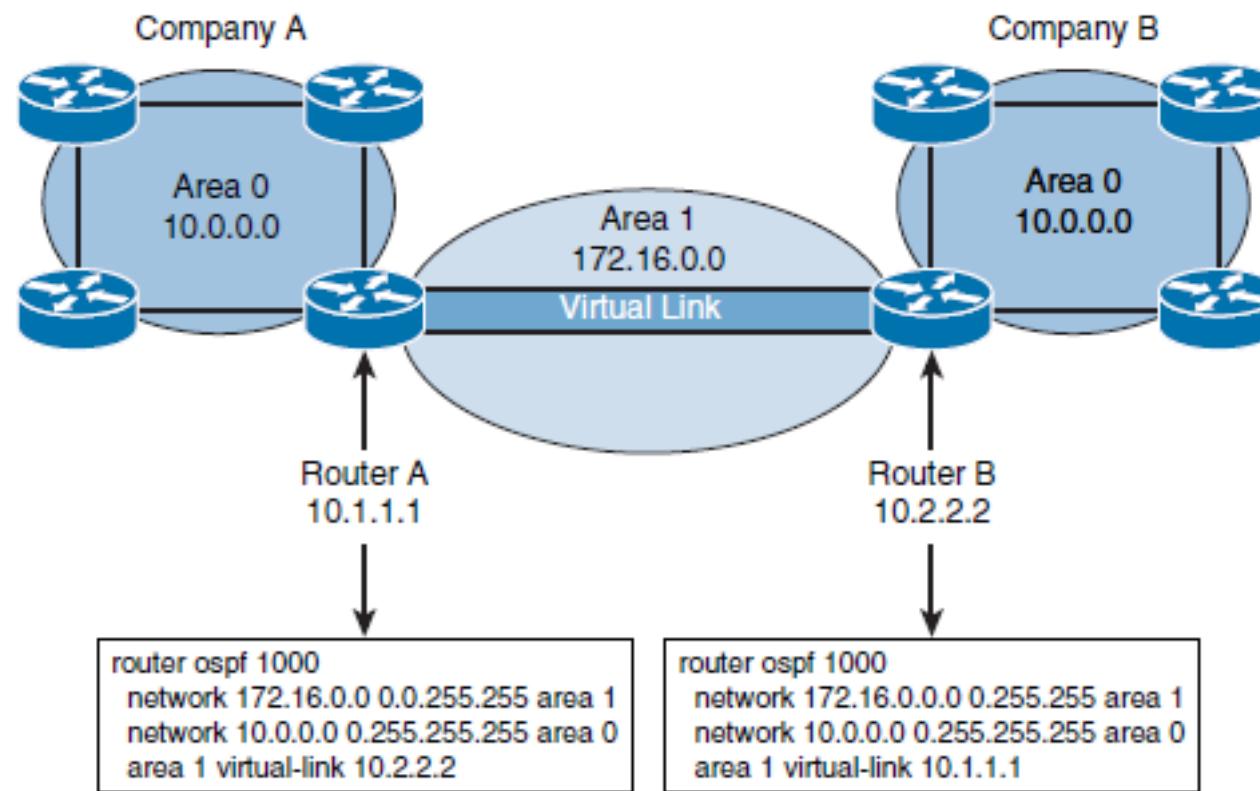


Virtual Links Configuration Issues

- The OSPF database treats the virtual link between ABR1 and ABR2 as a direct link.
- For greater stability, loopback interfaces are used as router IDs, and virtual links are created using these loopback addresses.
- The hello protocol works over virtual links as it does over standard links, in 10-second intervals.
- An LSA usually refreshes every 30 minutes. However, LSAs learned through a virtual link have the **DoNotAge** (DNA) option set so that the LSA does not age out. This DNA technique is required to prevent **excessive flooding** over the virtual link.



Configuring OSPF Virtual Links



```
area area-id virtual-link router-id [ authentication [ message-digest | null ][ hello-interval seconds ] [ retransmit-interval seconds ] [ transmitdelay seconds ] [ dead-interval seconds ] [[ authentication-key key ] | [ message-digest-key key-id md5 key ]]
```



Configuring OSPF Virtual Links

Parameter	Description
<i>area-id</i>	Specifies the area ID of the transit area for the virtual link. This ID can be either a decimal value or in dotted-decimal format, like a valid IP address. There is no default. The transit area cannot be a stub area.
<i>router-id</i>	Specifies the router ID of the virtual link neighbor. The router ID appears in the <i>show ip ospf</i> display. This value is in an IP address format. There is no default.
<i>authentication</i>	(Optional) Specifies an authentication type.
<i>message-digest</i>	(Optional) Specifies the use of MD5 authentication.
<i>null</i>	(Optional) Overrides simple password or MD5 authentication if configured for the area. No authentication is used.
<i>hello-interval seconds</i>	(Optional) Specifies the time (in seconds) between the hello packets that the Cisco IOS Software sends on an interface. The unsigned integer value is advertised in the Hello packets. The value must be the same for all routers and access servers attached to a common network. The default is 10 seconds.



Configuring OSPF Virtual Links

Parameter	Description
retransmit-interval <i>seconds</i>	(Optional) Specifies the time (in seconds) between LSA retransmissions for adjacencies belonging to the interface. The value must be greater than the expected round-trip delay between any two routers on the attached network. The default is 5 seconds.
transmit-delay <i>seconds</i>	(Optional) Specifies the estimated time (in seconds) to send an LSU packet on the interface. This integer value must be greater than 0. LSAs in the update packet have their age incremented by this amount before transmission. The default value is 1 second.
dead-interval <i>seconds</i>	(Optional) Specifies the time (in seconds) that must pass without hello packets being seen before a neighboring router declares the router down. This is an unsigned integer value. The default is 4 times the default hello interval, or 40 seconds. As with the hello interval, this value must be the same for all routers and access servers attached to a common network.
authentication-key <i>key</i>	(Optional) Specifies the password used by neighboring routers for simple password authentication. It is any continuous string of up to eight characters. There is no default value.
message-digest-key <i>key-id</i> md5 <i>key</i>	(Optional) Identifies the key ID and key (password) used between this router and neighboring routers for MD5 authentication. There is no default value.

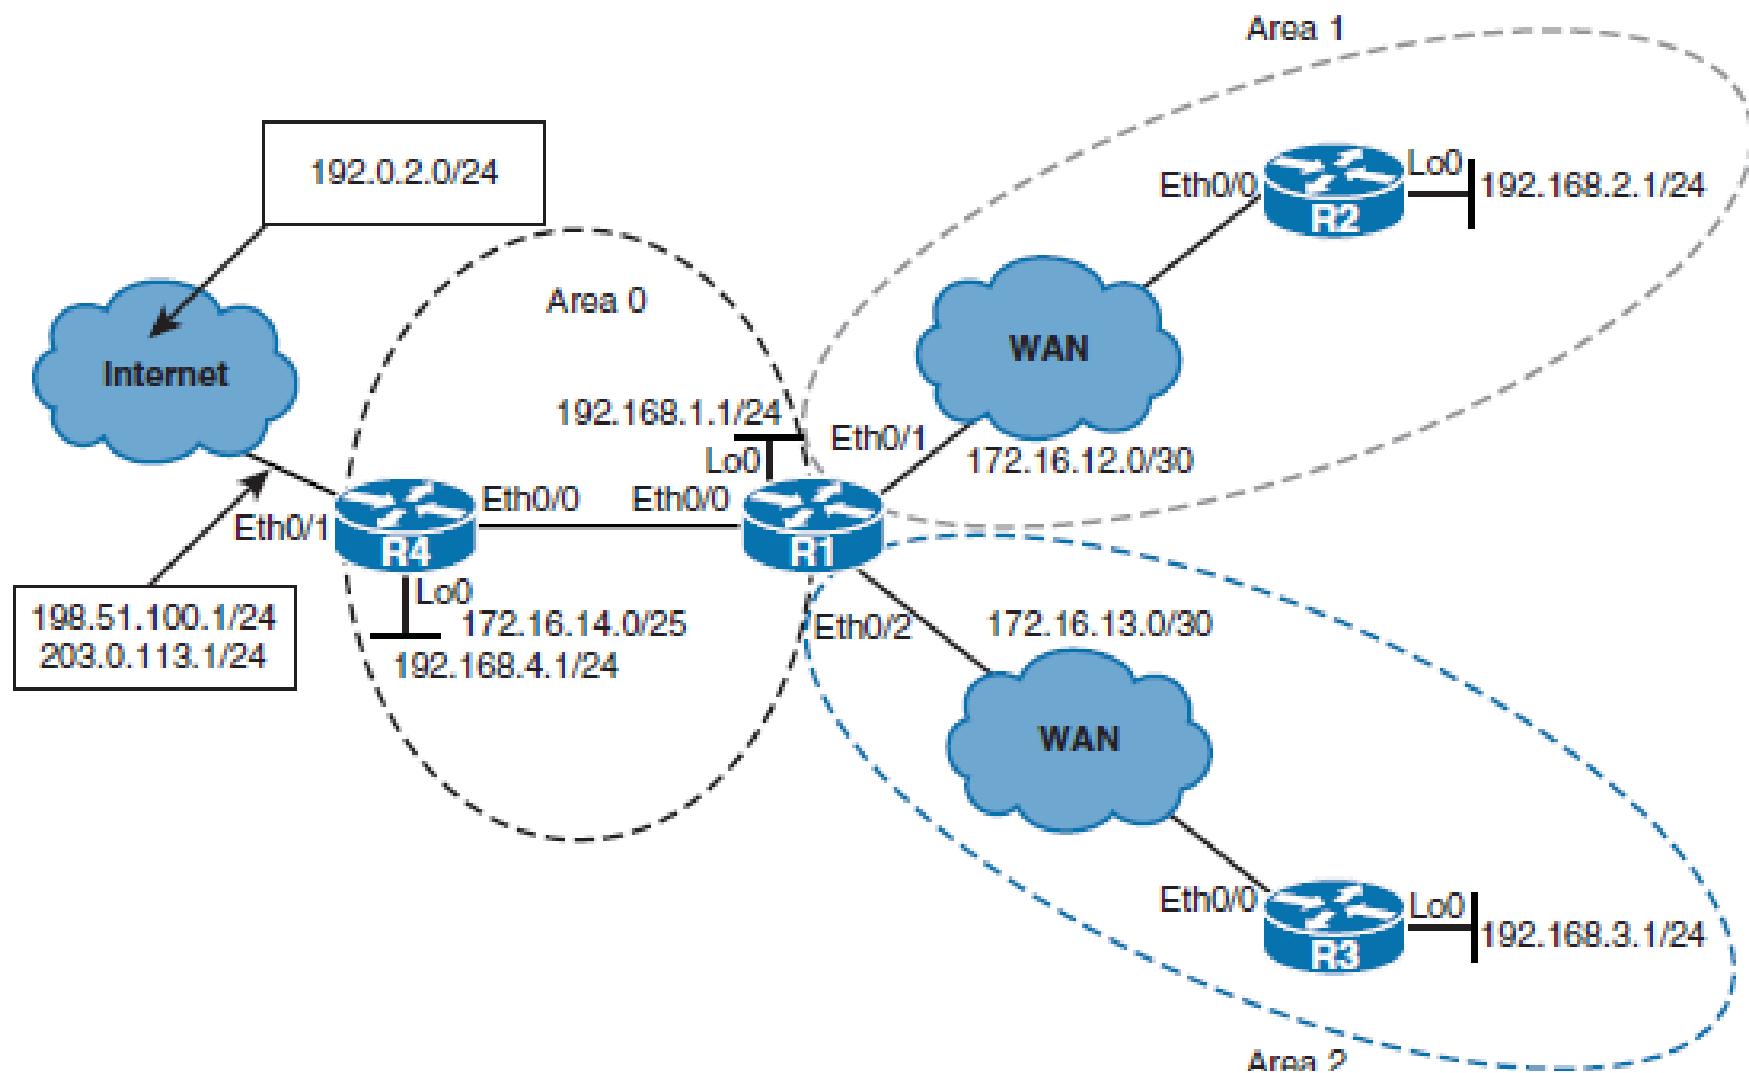


OSPF Stub Areas

- The **stub** and **totally stubby** areas are deployed to **reduce the size of the OSPF database and routing table**:
 - **Stub area:** This area type does **not accept** information about **routes external** to the autonomous system, such as routes from non-OSPF sources. If routers need to route to networks outside the autonomous system, they use a default route, indicated as 0.0.0.0. Stub areas **cannot contain ASBRs** (except that the ABRs may also be ASBRs). The stub area does not accept external routes.
 - **Totally stubby area:** This Cisco proprietary area type does not accept external autonomous system routes or summary routes from other areas internal to the autonomous system. If a router needs to send a packet to a network external to the area, it sends the packet using a default route. Totally stubby areas **cannot contain ASBRs** (except that the ABRs may also be ASBRs). A totally stubby area **does not accept external or interarea routes**.



Configuring OSPF Stub Areas





Configuring OSPF Stub Areas

```
R2# show ip route ospf
<Output omitted>

      172.16.0.0/16 is variably subnetted, 4 subnets, 3 masks
O IA    172.16.13.0/30 [110/20] via 172.16.12.1, 00:56:16, Ethernet0/0
O IA    172.16.14.0/25 [110/20] via 172.16.12.1, 00:56:16, Ethernet0/0
O IA    192.168.1.0/24 [110/11] via 172.16.12.1, 00:56:16, Ethernet0/0
O IA    192.168.3.0/24 [110/21] via 172.16.12.1, 00:54:50, Ethernet0/0
O IA    192.168.4.0/24 [110/21] via 172.16.12.1, 00:46:00, Ethernet0/0
O E2    198.51.100.0/24 [110/20] via 172.16.12.1, 00:01:47, Ethernet0/0
O E2    203.0.113.0/24 [110/20] via 172.16.12.1, 00:01:47, Ethernet0/0
```

```
R3# show ip route ospf
<Output omitted>

      172.16.0.0/16 is variably subnetted, 4 subnets, 3 masks
O IA    172.16.12.0/30 [110/20] via 172.16.13.1, 00:53:58, Ethernet0/0
O IA    172.16.14.0/25 [110/20] via 172.16.13.1, 00:53:58, Ethernet0/0
O IA    192.168.1.0/24 [110/11] via 172.16.13.1, 00:53:58, Ethernet0/0
O IA    192.168.2.0/24 [110/21] via 172.16.13.1, 00:53:58, Ethernet0/0
O IA    192.168.4.0/24 [110/21] via 172.16.13.1, 00:45:10, Ethernet0/0
O E2    198.51.100.0/24 [110/20] via 172.16.13.1, 00:00:57, Ethernet0/0
O E2    203.0.113.0/24 [110/20] via 172.16.13.1, 00:00:57, Ethernet0/0
```



Configuring OSPF Stub Areas

```
R1(config)# router ospf 1
R1(config-router)# area 1 stub
%OSPF-5-ADJCHG: Process 1, Nbr 2.2.2.2 on Ethernet0/1 from FULL to DOWN, Neighbor
Down: Adjacency forced to reset
```

```
R2(config)# router ospf 1
R2(config-router)# area 1 stub
%OSPF-5-ADJCHG: Process 1, Nbr 1.1.1.1 on Ethernet0/0 from LOADING to FULL, Loading
Done
```

- Configuring a stub area reduces the size of the LSDB inside the area, resulting in reduced memory requirements for routers in that area. External network LSAs (type 5), such as those that are redistributed from other routing protocols into OSPF, are not permitted to flood into a stub area.
- The **area stub** router configuration mode command is used to define an area as a stub area. Each router in the stub area must be configured with the **area stub** command. The
- Hello packets that are exchanged between OSPF routers contain a stub area flag that must match on neighboring routers. Until the **area 1 stub** command is enabled on R2 in this scenario, the adjacency between R1 and R2 will be down.



Configuring OSPF Stub Areas

```
R2# show ip route ospf
<Output omitted>

O*IA  0.0.0.0/0 [110/11] via 172.16.12.1, 00:19:27, Ethernet0/0
      172.16.0.0/16 is variably subnetted, 4 subnets, 3 masks
O  IA    172.16.13.0/30 [110/20] via 172.16.12.1, 00:19:27, Ethernet0/0
O  IA    172.16.14.0/25 [110/20] via 172.16.12.1, 00:19:27, Ethernet0/0
O  IA    192.168.1.0/24 [110/11] via 172.16.12.1, 00:19:27, Ethernet0/0
O  IA    192.168.3.0/24 [110/21] via 172.16.12.1, 00:19:27, Ethernet0/0
O  IA    192.168.4.0/24 [110/21] via 172.16.12.1, 00:19:27, Ethernet0/0

R2# ping 192.0.2.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.0.2.1, timeout is 2 seconds:
U.U.U
Success rate is 0 percent (0/5)

R2# ping 203.0.113.2
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 203.0.113.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms
```



Propagating a Default Route Using OSPF

```
R4# show ip route static
<Output omitted>
Gateway of last resort is 198.51.100.2 to network 0.0.0.0
S*    0.0.0.0/0 [1/0] via 198.51.100.2
R4(config)# router ospf 1
R4(config-router)# default-information originate
```



OSPF Totally Stubby Areas

```
R1(config)# router ospf 1
R1(config-router)# area 2 stub no-summary
%OSPF-5-ADJCHG: Process 1, Nbr 3.3.3.3 on Ethernet0/2 from FULL to
DOWN, Neighbor Down: Adjacency forced to reset
```

```
R3(config)# router ospf 1
R3(config-router)# area 2 stub
%OSPF-5-ADJCHG: Process 1, Nbr 1.1.1.1 on Ethernet0/0 from LOADING to FULL, Loading
Done
```



OSPF Totally Stubby Areas

```
R3# show ip route ospf
<Output omitted>
Gateway of last resort is 172.16.13.1 to network 0.0.0.0
O*IA  0.0.0.0/0 [110/11] via 172.16.13.1, 00:18:08, Ethernet0/0
```

```
R3# show ip ospf data
<Output omitted>
Summary Net Link States (Area 2)
```

Link ID	ADV Router	Age	Seq#	Checksum
0.0.0.0	1.1.1.1	1285	0x80000001	0x0093A6

```
R3# ping 192.0.2.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.0.2.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms
```



Cost of the Default Route in a Stub Area

- By default, the ABR of a stub area will advertise a **default route with a cost of 1**.
- You can change the cost of the default route by using the **area default-cost** command.
- The *default-cost* option provides the metric for the summary default route that is generated by the ABR into the stub area.

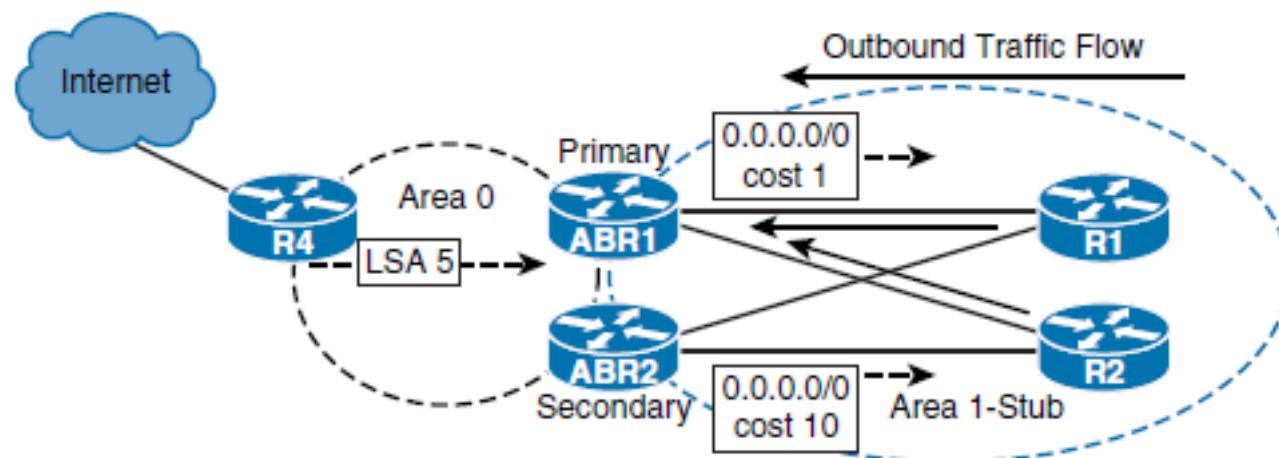
area *area-id* default-cost *cost*

Parameter	Description
<i>area-id</i>	Identifier for the stub or NSSA. The identifier can be specified as either a decimal value or as an IP address.
<i>cost</i>	Cost for the default summary route used for a stub or NSSA. The acceptable value is a 24-bit number.



Cost of the Default Route in a Stub Area

- The option of tuning the cost of the default route in the stub area is useful in stub areas with redundant exit points to the backbone area





The default-information originate Command

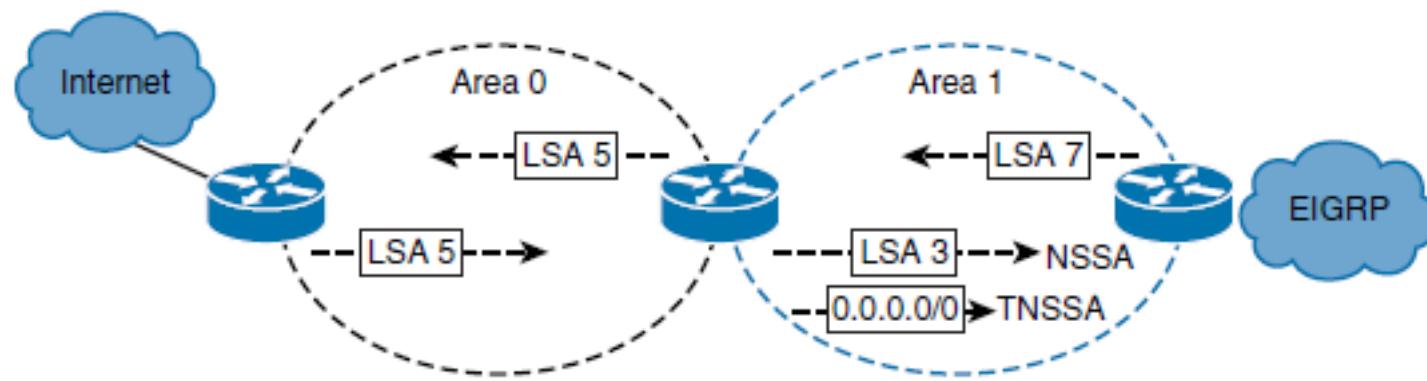
- **default-information originate [always] [metric *metric-value*] [metric-type *typevalue*] [route-map *map-name*]**

Parameter	Description
<i>always</i>	(Optional) Always advertises the default route regardless of whether the software has a default route.
<i>metric metric-value</i>	(Optional) Metric used for generating the default route. If you omit a value and do not specify a value using the default-metric router configuration command, the default metric value is 1. The value used is specific to the protocol.
<i>metric-type type-value</i>	(Optional) External link type associated with the default route advertised into the OSPF routing domain. It can be one of the following values: 1: Type 1 external route 2: Type 2 external route The default is type 2 external route.
<i>route-map map-name</i>	(Optional) Routing process will generate the default route if the route map is satisfied.



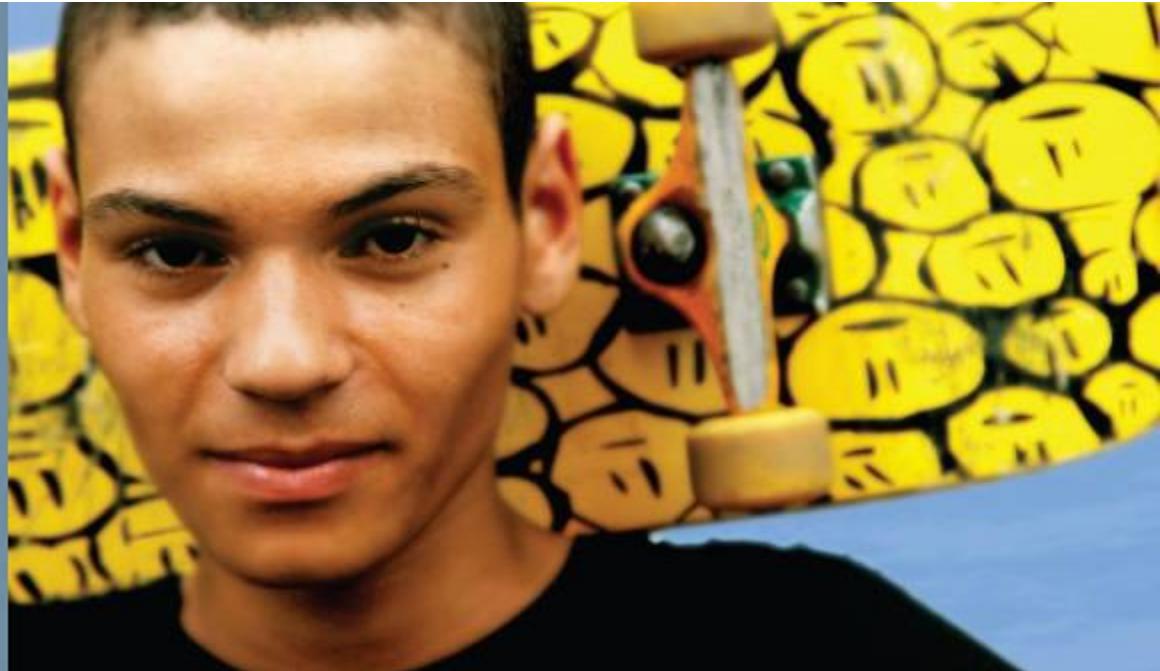
Other Stubby Area Types

- The NSSA is a nonproprietary extension of the existing stub area feature that allows the injection of external routes in a limited fashion into the stub area.



- The **area nssa** router configuration mode command is used to define each router in the NSSA area as not-so-stubby
- The **area nssa** command with the **no-summary** keyword is used to define the ABR as totally not-so-stubby.

OSPFv3



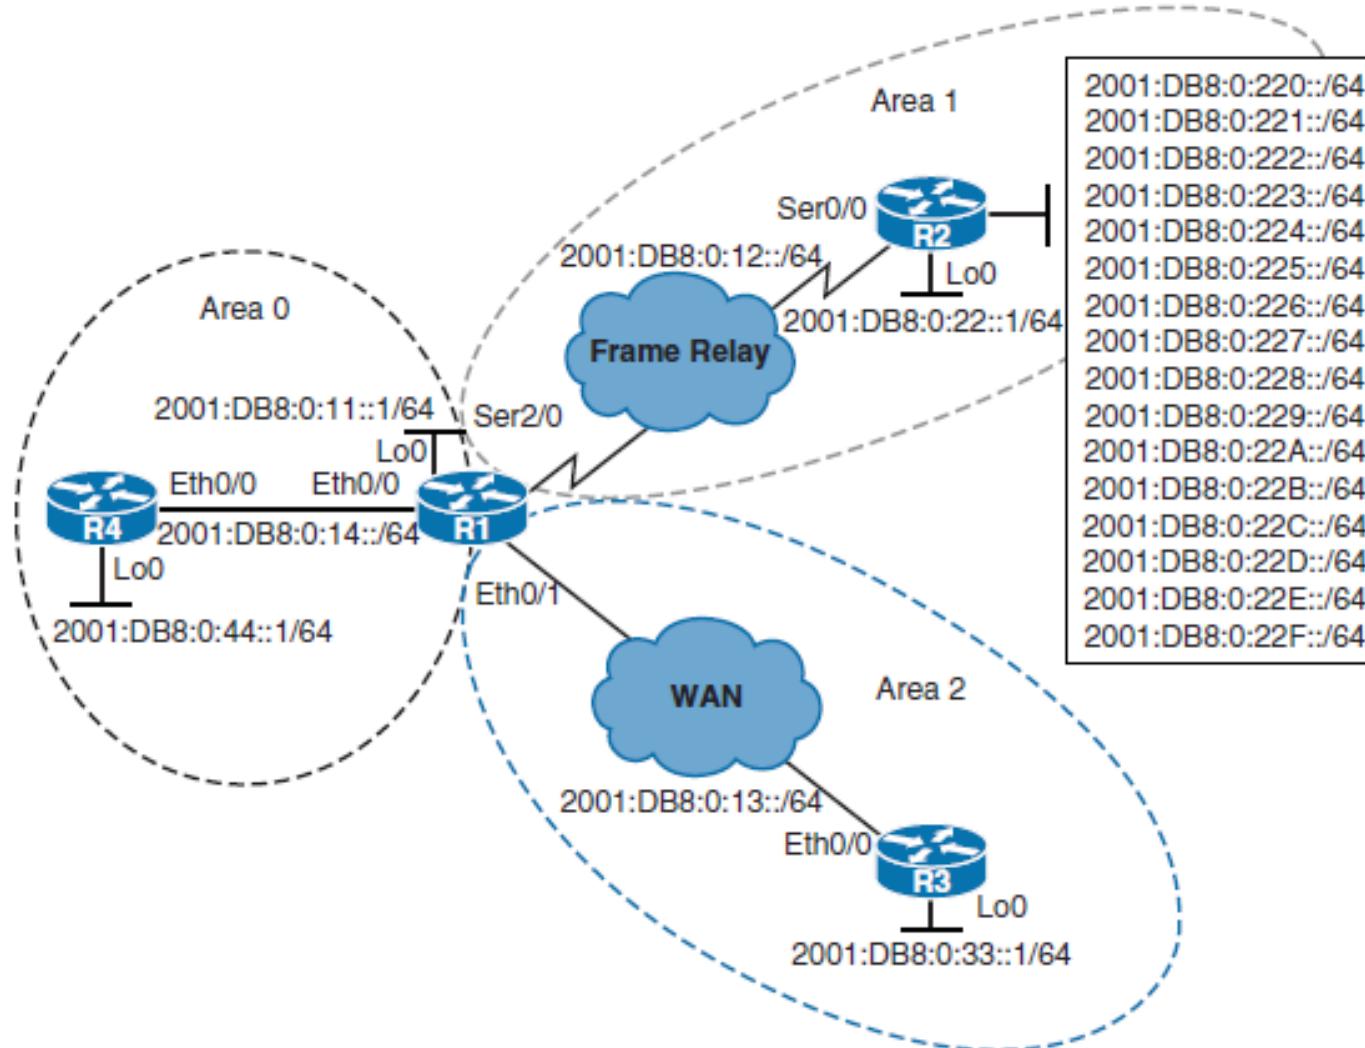


OSPFv3

- Implement OSPFv3 in a dual-stack (IPv4/IPv6) environment
- Configure external route summarization and load balancing in OSPFv3
- Explain the limitations and where you need to be careful when configuring OSPFv3



Configuring OSPFv3





Configuring OSPFv3

- Cisco IOS routers offer two OSPF configuration methods for IPv6:
- Using the traditional **ipv6 router ospf** global configuration command
- Using the new-style **router ospfv3** global configuration command



Configuring OSPFv3

- OSPFv3 is the IPv6-capable version of the OSPF routing protocol. It is a rewrite of the OSPF protocol to support IPv6, although the foundation remains the same as in IPv4 and OSPFv2.
- The OSPFv3 **metric is still based on interface cost**.
- The packet types and neighbor discovery **mechanisms are the same** in OSPFv3 as they are for OSPFv2.
- OSPFv3 also supports the same interface types, including broadcast, point-to-point, point-to-multipoint,
- NBMA, and virtual links.
- LSAs are still flooded throughout an OSPF domain, and many of the LSA types are the same, though a few have been renamed or newly created.



Enabling OSPFv3

```
R1(config)# ipv6 unicast-routing
R1(config)# ipv6 router ospf 1
R1(config-rtr)# router-id 1.1.1.1
R1(config-rtr)# passive-interface Loopback0
```

```
R1(config)# interface Loopback0
R1(config-if)# ipv6 ospf 1 area 0
R1(config-if)# exit
R1(config)# interface Ethernet0/0
R1(config-if)# ipv6 ospf 1 area 0
R1(config-if)# exit
R1(config)# interface Serial2/0
R1(config-if)# ipv6 ospf 1 area 1
R1(config-if)# exit
R1(config)# interface Ethernet0/1
R1(config-if)# ipv6 ospf 1 area 2
%OSPFv3-5-ADJCHG: Process 1, Nbr 4.4.4.4 on Ethernet0/0 from LOADING to FULL,
Loading Done
%OSPFv3-5-ADJCHG: Process 1, Nbr 3.3.3.3 on Ethernet0/1 from LOADING to FULL,
Loading Done
```



OSPFv3 Adjacencies and Routing Table

```
R1# show ipv6 ospf neighbor
```

```
OSPFv3 Router with ID (1.1.1.1) (Process ID 1)
```

Neighbor ID	Pri	State	Dead Time	Interface ID	Interface
4.4.4.4	1	FULL/DR	00:00:37	3	Ethernet0/0
3.3.3.3	1	FULL/DR	00:00:35	4	Ethernet0/1

```
R1# show ipv6 route ospf
```

```
<Output omitted>
```

```
O 2001:DB8:0:33::/64 [110/11]
    via FE80::A8BB:CCFF:FE00:AD10, Ethernet0/1
O 2001:DB8:0:44::/64 [110/11]
    via FE80::A8BB:CCFF:FE00:AE00, Ethernet0/0
```



Specifying the Neighbor on an NBMA Interface

- OSPF adjacencies over NBMA links require that IPv6 connectivity for both the link-local and the global addresses is established. Depending on the transport network, this may require mapping of IPv6 addresses to Layer 2 circuit identifiers

```
R1(config)# interface serial 2/0
R1(config-if)# ipv6 ospf neighbor FE80::2
%OSPFv3-5-ADJCHG: Process 1, Nbr 2.2.2.2 on Serial2/0 from LOADING to FULL, Loading
Done
```



Specifying the Neighbor on an NBMA Interface

```
R1# show running-config interface serial 2/0
Building configuration...

Current configuration : 404 bytes
!
interface Serial2/0
    ip address 172.16.12.1 255.255.255.252
    encapsulation frame-relay
    ipv6 address FE80::1 link-local
    ipv6 address 2001:DB8:0:12::1/64
    ipv6 ospf 1 area 1
    ipv6 ospf neighbor FE80::2
    serial restart-delay 0
    frame-relay map ip 172.16.12.2 102 broadcast
    frame-relay map ipv6 2001:DB8:0:12::2 102 broadcast
    frame-relay map ipv6 FE80::2 102 broadcast
    no frame-relay inverse-arp
end
```



OSPFv3 LSDB

```
R1# show ipv6 ospf database
```

OSPFv3 Router with ID (1.1.1.1) (Process ID 1)

Router Link States (Area 0)

ADV Router	Age	Seq#	Fragment ID	Link count	Bits
1.1.1.1	854	0x80000003	0	1	B
4.4.4.4	871	0x80000002	0	1	None

Net Link States (Area 0)

ADV Router	Age	Seq#	Link ID	Rtr count
4.4.4.4	871	0x80000001	3	2

Inter Area Prefix Link States (Area 0)

ADV Router	Age	Seq#	Prefix
1.1.1.1	845	0x80000001	2001:DB8:0:12::/64
1.1.1.1	845	0x80000001	2001:DB8:0:13::/64
1.1.1.1	845	0x80000001	2001:DB8:0:33::/64

Link (Type-8) Link States (Area 0)

ADV Router	Age	Seq#	Link ID	Interface
1.1.1.1	870	0x80000001	3	Ethernet0/0
4.4.4.4	1056	0x80000002	3	Ethernet0/0

Intra Area Prefix Link States (Area 0)

ADV Router	Age	Seq#	Link ID	Ref-ls type	Ref-LSID
1.1.1.1	865	0x80000003	0	0x2001	0
4.4.4.4	871	0x80000003	0	0x2001	0
4.4.4.4	871	0x80000001	3072	0x2002	3



Renamed LSA

- **Inter-area prefix LSAs for ABRs (Type 3):** Type 3 LSAs advertise internal networks to routers in other areas (inter-area routes). Type 3 LSAs may represent a single network or a set of networks summarized into one advertisement. Only ABRs generate summary LSAs. In OSPF for IPv6, addresses for these LSAs are expressed as prefix/ prefix length instead of address and mask. The default route is expressed as a prefix with length 0.
- **Inter-area router LSAs for ASBRs (Type 4):** Type 4 LSAs advertise the location of an ASBR. Routers that are trying to reach an external network use these advertisements to determine the best path to the next hop. ASBRs generate Type 4 LSAs.



New LSA

- **Link LSAs (Type 8):** Type 8 LSAs have local-link flooding scope and are never flooded beyond the link with which they are associated. Link LSAs provide the link-local address of the router to all other routers attached to the link. They inform other routers attached to the link of a list of IPv6 prefixes to associate with the link. In addition, they allow the router to assert a collection of option bits to associate with the network LSA that will be originated for the link.
- **Intra-area prefix LSAs (Type 9):** A router can originate multiple intra-area prefix LSAs for each router or transit network, each with a unique link-state ID. The linkstate ID for each intra-area prefix LSA describes its association to either the router LSA or the network LSA. The link-state ID also contains prefixes for stub and transit networks.



OSPFv3 for IPv4 and IPv6

- The newest OSPFv3 configuration approach utilizes a single OSPFv3 process.
- It is capable of **supporting IPv4 and IPv6 within a single OSPFv3 process.**
- OSPFv3 builds a single database with LSAs that carry IPv4 and IPv6 information.
- The OSPF adjacencies are established separately for each address family.
- Settings that are specific to an address family (IPv4/IPv6) are configured inside that address family router configuration mode.



Configuring OSPFv3 Using the router ospfv3

```
R1(config)# router ospfv3 1
R1(config-router)# router-id 1.1.1.1
R1(config-router)# passive-interface Loopback0
```

```
R1# show running-config | section router
router ospfv3 1
  router-id 1.1.1.1
  !
  address-family ipv6 unicast
    passive-interface Loopback0
    router-id 1.1.1.1
  exit-address-family
```

- The **address-family ipv6 unicast** has been automatically created on R1. Cisco IOS Software has parsed the previous old-style OSPFv3 configuration and found that the OSPF process was enabled only for IPv6.



OSPFv3 Old-Style OSPF Configuration Commands

```
interface Loopback0
    ipv6 ospf 1 area 0
!
interface Ethernet0/0
    ipv6 ospf 1 area 0
!
interface Ethernet0/1
    ipv6 ospf 1 area 2
!
interface Serial2/0
    ipv6 ospf 1 area 1
    ipv6 ospf neighbor FE80::2
```



OSPFv3 New-Style OSPF Configuration Commands

```
R1(config)# interface Loopback 0
R1(config-if)# ospfv3 1 ipv6 area 0
R1(config-if)# exit
R1(config)# interface Ethernet 0/0
R1(config-if)# ospfv3 1 ipv6 area 0
R1(config-if)# exit
R1(config)# interface Serial 2/0
R1(config-if)# ospfv3 1 ipv6 area 1
R1(config-if)# exit
R1(config)# interface Ethernet 0/1
R1(config-if)# ospfv3 1 ipv6 area 2
```

- The preferred interface mode command for the new style OSPFv3 configuration is the **ospfv3 process-id {ipv4|ipv6} area area-id** command.
- It allows you to selectively activate the OSPFv3 process for an address family (IPv4 or IPv6) on a given interface.



Enabling OSPFv3 for IPv4

```
R1(config)# interface Loopback0
R1(config-if)# ospfv3 1 ipv4 area 0
R1(config-if)# exit
R1(config)# interface Ethernet0/0
R1(config-if)# ospfv3 1 ipv4 area 0
R1(config-if)# exit
R1(config)# interface Ethernet0/1
R1(config-if)# ospfv3 1 ipv4 area 2
R1(config)# exit
R1(config-if)# interface Serial2/0
R1(config-if)# ospfv3 1 ipv4 area 1
R1(config-if)# ospfv3 1 ipv4 neighbor FE80::2
R1(config-if)# exit
R1(config)# router ospfv3 1
R1(config-router)# address-family ipv4 unicast
R1(config-router-af)# passive-interface Loopback0
%OSPFV3-5-ADJCHG: Process 1, IPv4, Nbr 0.0.0.0 on Serial2/0 from ATTEMPT to DOWN,
Neighbor Down: Interface down or detached
%OSPFV3-5-ADJCHG: Process 1, IPv4, Nbr 3.3.3.3 on Ethernet0/1 from LOADING to FULL,
Loading Done
%OSPFV3-5-ADJCHG: Process 1, IPv4, Nbr 4.4.4.4 on Ethernet0/0 from LOADING to FULL,
Loading Done
```



OSPFv3 Adjacencies for Both IPv4 and IPv6 Address Families

```
R1# show ospfv3 neighbor
```

```
OSPFv3 1 address-family ipv4 (router-id 1.1.1.1)
```

Neighbor ID	Pri	State	Dead Time	Interface ID	Interface
4.4.4.4	1	FULL/DR	00:00:34	3	Ethernet0/0
2.2.2.2	1	FULL/DR	00:01:38	3	Serial2/0
3.3.3.3	1	FULL/DR	00:00:36	4	Ethernet0/1

```
OSPFv3 1 address-family ipv6 (router-id 1.1.1.1)
```

Neighbor ID	Pri	State	Dead Time	Interface ID	Interface
4.4.4.4	1	FULL/DR	00:00:35	3	Ethernet0/0
2.2.2.2	1	FULL/DR	00:01:58	3	Serial2/0
3.3.3.3	1	FULL/DR	00:00:34	4	Ethernet0/1



OSPFv3 LSDB

```
R1# show ospfv3 database inter-area prefix
    OSPFv3 1 address-family ipv4 (router-id 1.1.1.1)
        Inter Area Prefix Link States (Area 0)
        LS Type: Inter Area Prefix Links
        Advertising Router: 1.1.1.1
        <Output omitted>
        Prefix Address: 172.16.12.0
        Prefix Length: 30, Options: None
        <Output omitted>

    OSPFv3 1 address-family ipv6 (router-id 1.1.1.1)
        Inter Area Prefix Link States (Area 0)
        LS Type: Inter Area Prefix Links
        Advertising Router: 1.1.1.1
        <Output omitted>
        Prefix Address: 2001:DB8:0:12::
        Prefix Length: 64, Options: None
        <Output omitted>
```



IPv4 Routing Table with OSPFv3 Routes

```
R1# show ip route ospfv3
<Output omitted>

      192.168.2.0/32 is subnetted, 1 subnets
O        192.168.2.2 [110/64] via 172.16.12.2, 00:27:49, Serial2/0
      192.168.3.0/32 is subnetted, 1 subnets
O        192.168.3.3 [110/10] via 172.16.13.2, 00:30:08, Ethernet0/1
      192.168.4.0/32 is subnetted, 1 subnets
O        192.168.4.4 [110/10] via 172.16.14.4, 00:30:08, Ethernet0/0
```

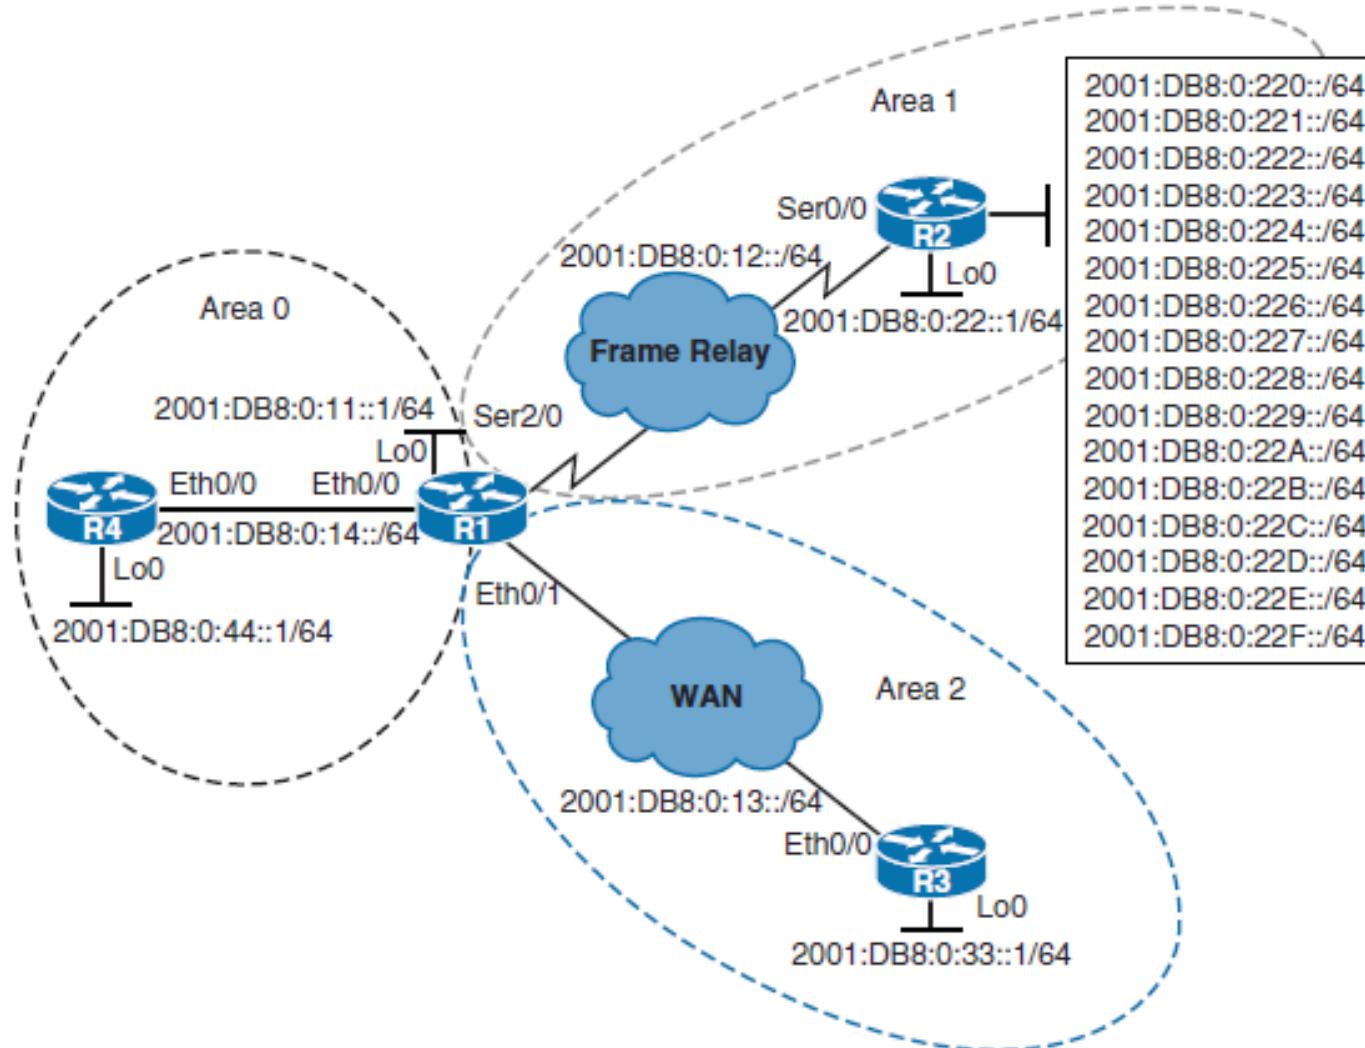


OSPFv3 ipv6 Routing Table

```
R3# show ipv6 route ospf
<Output omitted>
OI  2001:DB8:0:11::1/128 [110/10]
    via FE80::A8BB:CCFF:FE00:AB10, Ethernet0/1
OI  2001:DB8:0:12::/64 [110/74]
    via FE80::A8BB:CCFF:FE00:AB10, Ethernet0/1
OI  2001:DB8:0:14::/64 [110/20]
    via FE80::A8BB:CCFF:FE00:AB10, Ethernet0/1
OI  2001:DB8:0:22::/64 [110/75]
    via FE80::A8BB:CCFF:FE00:AB10, Ethernet0/1
OI  2001:DB8:0:44::/64 [110/21]
    via FE80::A8BB:CCFF:FE00:AB10, Ethernet0/1
OI  2001:DB8:0:220::/64 [110/75]
    via FE80::A8BB:CCFF:FE00:AB10, Ethernet0/1
OI  2001:DB8:0:221::/64 [110/75]
    via FE80::A8BB:CCFF:FE00:AB10, Ethernet0/1
OI  2001:DB8:0:222::/64 [110/75]
    via FE80::A8BB:CCFF:FE00:AB10, Ethernet0/1
```



Configuring OSPFv3





Area 2 Routers Configured as a Totally Stubby Area

```
R1(config)# router ospfv3 1
R1(config-router)# address-family ipv6 unicast
R1(config-router-af)# area 2 stub no-summary
%OSPFv3-5-ADJCHG: Process 1, IPv6, Nbr 3.3.3.3 on Ethernet0/1 from FULL to DOWN,
Neighbor Down: Adjacency forced to reset
```

```
R3(config)# router ospfv3 1
R3(config-router)# address-family ipv6 unicast
R3(config-router-af)# area 2 stub
%OSPFv3-5-ADJCHG: Process 1. IPv6. Nbr 1.1.1.1 on Ethernet0/1 from LOADING to FULL.
Loading Done
```

```
R3# show ipv6 route ospf
<Output omitted>
O I   ::/0 [110/11]
      via FE80::A8BB:CCFF:FE00:AB10, Ethernet0/1
R3# show ip route ospfv3
<Output omitted>
O IA     172.16.12.0/30 [110/74] via 172.16.13.1, 00:09:55, Ethernet0/1
O IA     172.16.14.0/25 [110/20] via 172.16.13.1, 00:09:55, Ethernet0/1
      192.168.1.0/32 is subnetted, 1 subnets
O IA     192.168.1.1 [110/10] via 172.16.13.1, 00:09:55, Ethernet0/1
      192.168.2.0/32 is subnetted, 1 subnets
O IA     192.168.2.2 [110/74] via 172.16.13.1, 00:09:55, Ethernet0/1
      192.168.4.0/32 is subnetted, 1 subnets
O IA     192.168.4.4 [110/20] via 172.16.13.1, 00:09:55, Ethernet0/1
```



OSPFv3 Summarization

```
R1(config)# router ospfv3 1
R1(config-router)# address-family ipv6 unicast
R1(config-router-af)# area 1 range 2001:DB8:0:220::/60
```

```
R4# show ipv6 route ospf
<Output omitted>
O  2001:DB8:0:11::1/128 [110/10]
    via FE80::A8BB:CCFF:FE00:AB00, Ethernet0/0
OI 2001:DB8:0:12::/64 [110/74]
    via FE80::A8BB:CCFF:FE00:AB00, Ethernet0/0
OI 2001:DB8:0:13::/64 [110/20]
    via FE80::A8BB:CCFF:FE00:AB00, Ethernet0/0
OI 2001:DB8:0:22::/64 [110/75]
    via FE80::A8BB:CCFF:FE00:AB00, Ethernet0/0
OI 2001:DB8:0:33::/64 [110/21]
    via FE80::A8BB:CCFF:FE00:AB00, Ethernet0/0
OI 2001:DB8:0:220::/60 [110/75]
    via FE80::A8BB:CCFF:FE00:AB00, Ethernet0/0
```



Configuring Advanced OSPFv3

- *Configuring the **summary-prefix** Command on an ASBR*

summary-prefix *prefix* [**not-advertise** | **tag tag-value**] [**nssa-only**]

Parameter	Description
<i>prefix</i>	IPv6 route prefix for the destination.
not-advertise	(Optional) Suppresses routes that match the specified prefix and mask pair. This keyword applies to OSPFv3 only.
tag tag-value	(Optional) Specifies the tag value that can be used as a match value for controlling redistribution via route maps. This keyword applies to OSPFv3 only.
nssa-only	(Optional) Limits the scope of the prefix to the area. Sets the NSSA-only attribute for the summary route (if any) generated for the specified prefix.

```
Router(config)# router ospfv3 1
Router(config-router)# address-family ipv6 unicast
Router(config-router-af)# summary-prefix 2001:db8:1::/56
```



OSPFv3 Caveats

- The OSPFv3 address families feature is supported as of Cisco IOS Release 15.1(3)S and Cisco IOS Release 15.2(1)T.
- Cisco devices that run software older than these releases and third-party devices will not form neighbor relationships with devices running the address family feature for the IPv4 address family because they do not set the address family bit.
- Therefore, those devices will not participate in the IPv4 address family SPF calculations and will not install the IPv4 OSPFv3 routes in the IPv6 Routing Information Base (RIB).



Chapter 3 Summary

- OSPF uses a two-layer hierarchical approach dividing networks into a backbone area (area 0) and nonbackbone areas.
- For its operation, OSPF uses five packet types: Hello, DBD, LSR, LSU, and LSAck.
- OSPF neighbors go through several different neighbor states before adjacency results in Full state.
- OSPF elects DR/BDR routers on a multiaccess segment to optimize exchange of information.
- The most common OSPF network types are point-to-point, broadcast, nonbroadcast, and loopback.
- OSPF uses several different LSA types to describe the network topology.
- LSAs are stored in an LSDB, which is synchronized with every network change.



Chapter 3 Summary

- OSPF calculates interface costs based on default reference bandwidth and interface bandwidth.
- Using SPF, OSPF determines the total lowest cost paths and selects them as the best routes.
- Intra-area routes are always preferred over interarea routes.
- Route summarization improves CPU utilization, reduces LSA flooding, and reduces routing table sizes.
- The **area range** command is used summarize at the ABR. The **summary-address** command is used to summarize at the ASBR.
- Default routes can be used in OSPF to prevent the need for specific route to each destination network.
- OSPF uses the **default-information originate** command to inject a default route.
- There are several OSPF area types: normal, backbone, stub, totally stubby, NSSA, and totally stubby NSSA.
- Use the **area area-id** command to define an area as stubby.
- Use the **area area-id stub** command with the **no-summary** keyword only on the ABR to define an area as totally stubby.



Chapter 3 Summary

- For stub areas, external routes are not visible in the routing table, but are accessible via the intra-area default route.
- For totally stubby areas, interarea and external routes are not visible in the routing table, but are accessible via the intra-area default route.
- OSPFv3 for IPv6 supports the same basic mechanisms that OSPFv2 for IPv4, including the use of areas to provide network segmentation and LSAs to exchange routing updates.
- OSPFv3 features two new LSA types and has renamed two traditional LSA types.
- OSPFv3 uses link-local addresses to source LSAs.
- OSPFv3 is enabled per-interface on Cisco routers.
- New-style OSPFv3 and traditional OSPFv3 for IPv6, configured with **ipv6 router ospf**, can coexist in the network to provide IPv6 routing.



Chapter 3 Labs

- **CCNPv7_ROUTE_Lab3-1 OSPF-Virtual-Links**
- **CCNPv7_ROUTE_Lab3-2 Multi-Area-OSPF**
- **CCNPv7_ROUTE_Lab3-3 OSPFv3-Address-Families**

Cisco | Networking Academy®

Mind Wide Open™



Acknowledgment

- Some of images and texts are from Implementing Cisco IP Routing (ROUTE) Foundation Learning Guide by Diane Teare, Bob Vachon and Rick Graziani (1587204568)
- Copyright © 2015 – 2016 Cisco Systems, Inc.
- Special Thanks to *Bruno Silva*